

Advanced Data Computing (405 @ LU) Notes

Abhishek Chakraborty

August 20, 2026

Table of contents

Welcome	6
I SQL	7
Introduction: Beyond the Flat File	8
The Relational Revolution	8
SQL as a Declarative Language	9
SQL vs. R	10
Industry Ubiquity	10
Cautionary Note	11
Further Resources	11
1 Accessing Databases and Database Tables	12
1.1 load packages	12
1.2 databases and accessing them	12
1.3 database tables and accessing them	13
1.4 terminate connection	20
2 Writing SQL Queries	21
2.1 load packages and connect to database server	21
2.2 SQL in R: Part 1 - the 'dbplyr' R package will directly translate 'dplyr' code into SQL	22
2.3 SQL in R: Part 2 - using the DBI R package to write SQL queries in r code chunks	26
2.4 SQL in R: Part 3 - writing SQL queries in sql code chunks	27
2.4.1 query 1	31
2.4.2 query 2	34
2.4.3 query 3	35
2.4.4 query 4	36
2.4.5 query 5	41
2.4.6 query 6	42
2.4.7 query 7	43
2.4.8 query 8	44
2.5 calling SQL generated datasets into R for data viz	46
2.5.1 Method 1: using DBI::dbGetQuery()	46
2.5.2 Method 2: using dbplyr with raw SQL	46

2.5.3	Method 3: Quarto chunk options	46
2.6	terminate connection	46
3	Joins, Unions, and Subqueries	47
3.1	load packages and connect to database server	47
3.1.1	flights table	48
3.1.2	carriers table	50
3.1.3	airports table	50
3.2	Joins (Horizontal Integration)	52
3.2.1	INNER JOIN	52
3.2.2	LEFT JOIN	54
3.3	Unions (Vertical Integration)	56
3.4	Subqueries	57
3.5	terminate connection	58
4	Creating Databases	59
4.1	load packages and connect to database server	59
4.2	accessing data (csv files) from GitHub (or, elsewhere)	60
4.3	data types in SQL	62
4.4	importing data	62
4.5	checks	63
4.6	shortcut - a fast (but fragile) route	65
4.7	terminate connection	67
II	Text Analysis	68
5	Case Study of Shakespeare's Macbeth using Regular Expressions	69
5.1	resources	69
5.2	packages	69
5.3	data	69
5.4	strings and regular expressions	70
5.5	about regular expressions grammar	70
5.5.1	metacharacter	70
5.5.2	character sets	71
5.5.3	alternation	71
5.5.4	anchors (power-dollar)	71
5.5.5	repetitions	71
5.6	a case study - life and death in Macbeth and speaking patterns	72
6	Case Study of Donald Trump's tweets during the 2016 US presidential election	74
6.1	resources	74
6.2	packages	74

6.3	data	74
6.4	some initial exploration	75
6.5	tokens	76
6.6	stop words	76
6.7	further exploration	77
6.8	sentiment analysis	78
7	Case Study of Data Science Papers from arXiv.org	81
7.1	resources	81
7.2	packages	81
7.3	data	81
7.4	initial data exploration and cleaning	82
7.5	corpora and tokenization	82
7.6	wordcloud	83
7.7	bigrams	83
7.8	tf-idf and document term matrices	84
III	Geospatial Data	86
8	Recreating John Snow's 1854 map of the Broad Street cholera outbreak.	87
8.1	resources	87
8.2	packages	87
8.3	data	87
8.4	projections	88
8.5	CRS - coordinate reference system	88
8.6	what about our dataset?	89
IV	Code Performance	91
9	Code Performance	92
9.1	resources	92
9.2	code profiling	92
9.3	improving performance with Rcpp	93
9.4	parallelization with future and furrr	94
9.5	efficient file handling	95
V	R Package	97
10	Building an R Package	98
10.1	check for available package names	98
10.2	create package	98

10.3	create functions within the package	99
10.4	add license	100
10.5	add documentations	101
10.6	add assertions	101
10.7	unit testing	101
10.8	add dataset	102
10.9	add README	103
10.10	add/build vignettes	103
10.11	build package website	104

Welcome

These notes have been developed for the CMSC/STAT 405: Advanced Data Computing course at [Lawrence University](#), Appleton, Wisconsin. This course is intended for upper-level undergraduate students who have previously taken a course on R-programming (that covers the `tidyverse` ecosystem, webscraping, and Shiny apps).

We use R (except while writing SQL code), and access it through Posit Cloud, as our primary computing environment. These notes have been heavily adopted and adapted from the following excellent open-source resources:

- [R for Data Science \(2e\)](#) by Hadley Wickham, Mine Çetinkaya-Rundel, and Garrett Grolemund.
- [Advanced R](#) by Hadley Wickham .
- [Modern Data Science with R \(3e\)](#) by Benjamin Baumer, Daniel Kaplan, and Nicholas Horton.
- [Text Mining with R](#) by Julia Silge and David Robinson.
- [R Packages \(2e\)](#) by Hadley Wickham and Jennifer Bryan.

Besides these wonderful texts, some additional resources are mentioned at the beginning of each module and/or chapter. Some of the prose have been written in collaboration with AI. Reader comments and feedback are greatly appreciated. To report errors or bugs, please post an issue or raise a pull request [here](#).

Part I

SQL

Introduction: Beyond the Flat File

In most introductory statistical and data science coursework (such as CMSC/STAT 205: Data-Scientific Programming), we typically work with *static* data, that is, a (or, a few) `.csv` or `.xlsx` file(s), which we can load entirely into our computer's active memory (RAM) - this is similar to loading dataset(s) into the **Environment** pane in the top right of our Posit Cloud screen.

While this setup is highly effective for isolated research questions and pedagogical exercises inside the classroom, data in the real-world, or simply, datasets, are rarely static, and they are almost always larger than what can fit into a single machine's memory. This is common in most industry environments and complex organizational structures.

This is where **Database Management Systems (DBMS)** become necessary. A DBMS is a software system designed to store data in a structured format, and which can also be used to retrieve, and manage data efficiently and securely. Moving from flat (static) files to a robust database framework solves three fundamental infrastructural problems:

1. **Scale:** Databases are optimized to query gigabytes or terabytes of data dynamically without requiring the entire dataset to be loaded into our local environment.
2. **Concurrency:** A flat file cannot easily be read and written to by hundreds of users simultaneously without data corruption. Databases handle concurrent actions safely, that is, allow for multiple users to access the data simultaneously, maintaining a single source of truth.
3. **Integrity:** Databases enforce strict rules about what data can enter the system (schemas), ensuring that a numeric column cannot accidentally ingest a string of text.

In a database, data are stored in the form of tables. Conceptually, a table in a database is no different from the data frames which are used in R. They have the same rectangular format with horizontal rows and vertical columns. However, there are two major differences between a data frame that we are used to (from 205) and a table in a database.

1. **Size:** Tables in a database are too large to be stored in our computer's memory unlike typical data frames we used previously. Thus, we need to store tables remotely on a hard drive outside of our personal computers. This, in turn, make the data secure and can be accessed by multiple users at a time.
2. **Relationality:** Tables in a database are usually linked with a **key**.

The Relational Revolution

Prior to the 1970s, data storage was highly rigid. Systems used hierarchical or network models, meaning that to extract data, a programmer had to write complex procedural code detailing

the exact physical navigation path through the computer's storage layer. If the physical storage hardware changed, the software broke.

In 1970, Edgar F. Codd, an IBM computer scientist, published a landmark paper: [*A Relational Model of Data for Large Shared Data Banks*](#). Codd proposed separating the physical storage of data from its logical organization. Data would be organized into relations (tables) consisting of tuples (rows) and attributes (columns). Connections between data were established through shared keys rather than physical pointers.

To interact with this new relational model, IBM developed **SEQUEL** (Structured English Query Language) in the mid-1970s, which was later shortened to **SQL** (Structured Query Language). SQL became an ANSI standard in 1986 and fundamentally changed how humans interact with data systems.

To digest the concept of relational databases, let's think of this setup. Suppose Lawrence has been storing the records of each student from every graduating class. So, unsurprisingly, that's a **lot** of data. Now, think about how do we want to store these data - say we store students names, their IDs, date of births, classes taken every term, etc. Do we want to store all these data into a single file? Definitely not, for a variety of reasons. Firstly, we would want to separate certain pieces of information. Data about names, date of births, IDs, for example, are unlikely to change. We would like to keep them in one table. Data about courses can be kept in a separate table. These two tables would be related by a key. Now, do we want the key to be student names or their IDs? Which one is more efficient to store in terms of individual entries? Secondly, storing data in the above way lets us update only one table (containing information about courses) every term without touching the other table (student demographics which are unlikely to change). That is more efficient and secure to manage as well. Thirdly, what happens if a student does change their name? Then we can update only one entry in one table, rather than updating multiple entries (if demographic data and course-specific data were stored in a combined table).

SQL as a Declarative Language

Most general-purpose programming languages we encounter are *imperative*. We write sequential instructions telling the computer exactly *how* to achieve a result.

SQL is fundamentally different; it is a *declarative* language. When we write SQL, we describe *what* result we want, and the database engine's query optimizer decides the most efficient algorithmic path to retrieve it. We do not write loops or dictate the order of operations. We state our logical requirements, and the engine handles the execution plan.

This abstraction makes SQL incredibly powerful. A query written to extract ten rows from a table of a thousand records uses the exact same syntax as a query operating on a table of ten billion records distributed across a cloud cluster.

SQL is a query language meant for working with relational databases. There exists many dialects of SQL (just like a regular spoken language). Translating between the dialects is not always easy although once we learn how to program in one dialect, we will be able to (hopefully) pick up any of the other SQL dialects. We will use the MySQL dialect in the upcoming chapters, unless mentioned otherwise. MySQL is among the most popular implementations of SQL and it is open source. MySQL is based on a **client-server model**, meaning, the data live on a powerful computer (the server) and we connect to the data from our own computer (the client).

SQL vs. R

R is a statistical computing environment that is developed for the purpose of data analysis. SQL is designed as a robust solution for data management and extraction, its analytic capabilities are very limited. Although, we will notice that there are similarities in the ways we think about and write R and SQL codes, they serve fundamentally different purposes. In the next chapters, we will alternate between SQL and R code, and use them for their specific tasks (for example, we will extract data from a database using SQL and visualize them using R). It's important for a data scientist to have both languages in their toolkit.

Industry Ubiquity

Despite being decades old, SQL is the undisputed *lingua franca* of modern data work. While specialized frameworks rise and fall, SQL remains the underlying connective tissue for data engineering, analytics, and applied statistics.

In contemporary industry contexts, SQL is used to:

- Extract subsets of data from enterprise data warehouses (like Snowflake, Amazon Redshift, or Google BigQuery) before bringing the manageable subset into an analytical environment for modeling.
- Define metrics and dashboards in business intelligence tools.
- Perform robust, in-database feature engineering prior to training predictive models.

Mastering SQL is not just about learning a syntax; it is about learning how to think in sets and relations. It bridges the gap between ad-hoc statistical analysis and production-grade data infrastructure.

Cautionary Note

In the next four chapters, we will use both R and SQL code chunks. Shown below are an example of each; notice (1) the header which defines the language (although this would be invisible later on), and (2) the difference in how to comment out code, that should be your cue to identify the specific code chunk.

💡 R Code Chunk

```
```{r}

your R code below

```
```

💡 SQL Code Chunk

```
```{sql}
#| connection: connection_to_your_server

-- your SQL code below

```
```

Further Resources

Besides the resources mentioned at the beginning of the text, I suggest you check out

- [Data Science, the SQL](#) notes by Jo Hardin.
- [Notes from Database Systems at UC Berkeley](#)
- [Introduction to dbplyr](#)
- [Hands-on Database \(2e\) - Steve Conger](#)
- [SQL in a Nutshell](#) by Kline et al.
- Documentations for different SQL dialects - [MySQL](#), [PostgreSQL](#), and [SQLite](#)

1 Accessing Databases and Database Tables

1.1 load packages

To interact with a database using R, we need a specific set of tools. The `tidyverse` collection of R packages provides our standard data manipulation framework, while DBI (Database Interface) serves as the common architecture for communicating with **Relational DataBase Management Systems (RDBMS)**. Because different databases (PostgreSQL, MySQL, SQLite) speak slightly different dialects of SQL, we also need a specific driver. In this case, we use `RMariaDB` to communicate with MySQL/MariaDB servers. Finally, `dbplyr` is the critical translation layer that allows us to write `dbplyr` code that R automatically converts into SQL.

```
library(tidyverse)
library(mdsr)
library(DBI)
library(RMariaDB)
library(dbplyr)
```

1.2 databases and accessing them

Relational databases operate on a **client-server** model. In this architecture, the data does not live on our local machine. Instead, it is housed on a centralized remote server. Our R session (the client) must send a formal request over a network to access that data.

We initialize this connection using the `mdsr` package, which provides a convenient wrapper function to connect to a public, read-only educational database.

Here, `con_air` is the connection to the `airlines` database on the `scidb` server located on Amazon Web Services. Check out more details on the help page (`?mdsr::dbConnect_scidb`).

```
# connecting to the airlines database

con_air <- mdsr::dbConnect_scidb("airlines")

# remember the name of the connection
```

The Client-Server Analogy

Think of the **client-server** model like a restaurant. We (the client) sit at a table and look at the menu. We do not go into the kitchen to cook the food ourselves. Instead, we send our order (the query) via a waiter (the connection) to the kitchen (the server). The kitchen processes the heavy lifting and sends back only our specific meal (the results).

While wrapper functions like the one above are convenient for teaching, they hide the underlying mechanics. In an industry setting, connecting to a database requires explicit routing and authentication parameters. We must specify the exact server location, the name of the database, our username, and our password.

Here is what an explicit connection via DBI looks like:

```
# access wai database

con_wai <- DBI::dbConnect(
  RMariaDB::MariaDB(),
  dbname = "wai",
  host = "scidb.smith.edu",
  user = "waiuser",
  password = "smith_waiDB")
```

Password Security

It is not ideal and hence, not recommended, to share passwords. The code chunk above is shown only for teaching purposes.

1.3 database tables and accessing them

Once a connection is established, the first step is typically reconnaissance: we need to know what tables exist within the database.

We can achieve this using R syntax via the DBI package:

```
# lists tables in the airlines database

DBI::dbListTables(con_air)
```

```
[1] "airports"      "planes"        "carriers"      "flights_summary"
[5] "flights"
```

Alternatively, because Quarto is a polyglot authoring framework, we can write raw SQL directly. By specifying `connection = "con_air"` or using `#| connection: con_air` in the chunk options, Quarto knows to pass the code inside the chunk directly to the database rather than to the R interpreter.

Note, the following is a SQL code chunk.

```
-- lists tables in the database

SHOW TABLES;
```

Table 1.1: 5 records

| Tables_in_airlines |
|--------------------|
| airports |
| carriers |
| flights |
| flights_summary |
| planes |

i SQL Syntax Conventions

Notice that we wrote `SHOW TABLES` in all capital letters and ended the statement with a semicolon (`;`). While SQL is technically case-insensitive, capitalizing SQL keywords is a universal standard that drastically improves readability, distinguishing commands from table and column names. The semicolon acts as a terminator, signaling to the server that the query is complete and ready for execution.

Knowing the tables exist is not enough; we must establish a reference to a specific table to query it. We do this using `dplyr::tbl()`. This creates a lazy reference to a remote database table without downloading it. Operations on this object are translated to SQL and executed on the server only when results are actually needed.

Here, we query the `flights` table from the established connection to the `airlines` database.

```
# reference to flights table from the airlines database

flights <- dplyr::tbl(con_air, "flights")
```

i Object Distinctions

It's easy to conflate the objects we have created thus far. Notice the distinction between `con_air` (the connection object), `airlines` (the database), and `flights` (the table reference). The following hierarchy would be helpful

server → database → table → table reference

Notice the output of `class(flights)`. It is not a standard `data.frame` or `tbl_df`. It is a `tbl_dbi` and `tbl_sql`. This distinction is profound because it introduces the concept of lazy evaluation.

```
# inspect the object type
```

```
class(flights)
```

```
[1] "tbl_MariaDBConnection" "tbl_dbi"          "tbl_sql"
[4] "tbl_lazy"             "tbl"
```

When we print the `flights` object below, it provides a quick glimpse of the columns and the first few rows. However, we will notice it says ?? for the total number of rows. This is because the data is not actually in R's memory. R has merely established a pointer to the table on the server. Because counting tens of millions of rows takes computational time, the database engine skips that step until explicitly commanded otherwise. It is “**lazy**”—it does as little work as possible.

```
# attempting to print the flights object
```

```
flights
```

```
# A query: ?? x 21
```

```
# Database: mysql [mdsr_public@mdsr.crcbo51tmesf.us-east-2.rds.amazonaws.com:3306/airlines]
```

| | year | month | day | dep_time | sched_dep_time | dep_delay | arr_time | sched_arr_time |
|---|-------|-------|-------|----------|----------------|-----------|----------|----------------|
| | <int> | <int> | <int> | <int> | <int> | <int> | <int> | <int> |
| 1 | 2013 | 10 | 1 | 2 | 10 | -8 | 453 | 505 |
| 2 | 2013 | 10 | 1 | 4 | 2359 | 5 | 730 | 729 |
| 3 | 2013 | 10 | 1 | 11 | 15 | -4 | 528 | 530 |
| 4 | 2013 | 10 | 1 | 14 | 2355 | 19 | 544 | 540 |
| 5 | 2013 | 10 | 1 | 16 | 17 | -1 | 515 | 525 |
| 6 | 2013 | 10 | 1 | 22 | 20 | 2 | 552 | 554 |
| 7 | 2013 | 10 | 1 | 29 | 35 | -6 | 808 | 816 |

```

8 2013    10    1    29          35      -6    449          458
9 2013    10    1    31          30       1    519          538
10 2013   10    1    32          33      -1    557          606
# i more rows
# i 13 more variables: arr_delay <int>, carrier <chr>, tailnum <chr>,
#   flight <int>, origin <chr>, dest <chr>, air_time <int>, distance <int>,
#   cancelled <int>, diverted <int>, hour <int>, minute <int>, time_hour <dtm>

```

i Lazy tbl

Notice that typical exploratory functions such as `str` and `glimpse` would work differently with the `flights` object compared to datasets loaded in the local session memory.

```
str(flights)
```

```

List of 3
 $ con      :Formal class 'MariaDBConnection' [package "RMariaDB"] with 7 slots
   .. ..@ ptr          :<pointer: 0x57311fa6b210>
   .. ..@ host          : chr "mdsr.crcbo51tmesf.us-east-2.rds.amazonaws.com"
   .. ..@ db            : chr "airlines"
   .. ..@ load_data_local_infile: logi FALSE
   .. ..@ bigint        : chr "integer64"
   .. ..@ timezone      : chr "UTC"
   .. ..@ timezone_out   : chr "UTC"
 $ src      :List of 2
   ..$ con   :Formal class 'MariaDBConnection' [package "RMariaDB"] with 7 slots
     .. .. ..@ ptr          :<pointer: 0x57311fa6b210>
     .. .. ..@ host          : chr "mdsr.crcbo51tmesf.us-east-2.rds.amazonaws.com"
     .. .. ..@ db            : chr "airlines"
     .. .. ..@ load_data_local_infile: logi FALSE
     .. .. ..@ bigint        : chr "integer64"
     .. .. ..@ timezone      : chr "UTC"
     .. .. ..@ timezone_out   : chr "UTC"
   ..$ disco: NULL
   ..- attr(*, "class")= chr [1:4] "src_MariaDBConnection" "src_dbi" "src_sql" "src"
 $ lazy_query:List of 5
   ..$ x      : 'dbplyr_table_path' chr "`flights`"
   ..$ vars    : chr [1:21] "year" "month" "day" "dep_time" ...
   ..$ group_vars: chr(0)
   ..$ order_vars: NULL

```



```

..$ frame      : NULL
..- attr(*, "class")= chr [1:3] "lazy_base_remote_query" "lazy_base_query" "lazy_query"
- attr(*, "class")= chr [1:5] "tbl_MariaDBConnection" "tbl_dbi" "tbl_sql" "tbl_lazy" ...

```

```
glimpse(flights) # notice the ?? on the number of rows
```

```

Rows: ??
Columns: 21
$ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2~
$ month     <int> 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, ~
$ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
$ dep_time  <int> 2, 4, 11, 14, 16, 22, 29, 29, 31, 32, 33, 34, 38, 39, 4~
$ sched_dep_time <int> 10, 2359, 15, 2355, 17, 20, 35, 35, 30, 33, 35, 2328, 4~
$ dep_delay <int> -8, 5, -4, 19, -1, 2, -6, -6, 1, -1, -2, 66, -2, -~
1, -5~
$ arr_time  <int> 453, 730, 528, 544, 515, 552, 808, 449, 519, 557, 445, ~
$ sched_arr_time <int> 505, 729, 530, 540, 525, 554, 816, 458, 538, 606, 457, ~
$ arr_delay <int> -12, 1, -2, 4, -10, -2, -8, -9, -19, -9, -~
12, 59, -7, --
$ carrier   <chr> "AA", "FL", "AA", "AA", "UA", "UA", "US", "AS", "DL", "~
$ tailnum   <chr> "N201AA", "N344AT", "N3KMAA", "N3ENAA", "N38473", "N458~
$ flight    <int> 2400, 710, 1052, 2392, 1614, 291, 436, 108, 1896, 687, ~
$ origin    <chr> "LAX", "SFO", "SFO", "SEA", "LAX", "SFO", "LAX", "ANC",~
$ dest      <chr> "DFW", "ATL", "DFW", "ORD", "IAH", "IAH", "CLT", "SEA",~
$ air_time  <int> 149, 247, 182, 191, 157, 188, 256, 181, 147, 183, 177, ~
$ distance  <int> 1235, 2139, 1464, 1721, 1379, 1635, 2125, 1448, 1399, 1~
$ cancelled <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ diverted  <int> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ hour      <int> 0, 23, 0, 23, 0, 0, 0, 0, 0, 0, 0, 23, 0, 0, 0, 0, 1, 0~
$ minute    <int> 10, 59, 15, 55, 17, 20, 35, 35, 30, 33, 35, 28, 40, 40,~
$ time_hour <dtm> 2013-10-01 00:10:00, 2013-10-01 23:59:00, 2013-~
10-01 0~

```

To understand the architecture of the table without pulling data, we can use the `DESCRIBE` command in `SQL`. This returns column/field names, data types, and constraints for a table. This is a recommended first step before working with the table.

```
DESCRIBE flights;
```

Table 1.2: Displaying records 1 - 10

| Field | Type | Null | Key | Default | Extra |
|----------------|-------------|------|-----|---------|-------|
| year | smallint(4) | YES | | NA | |
| month | smallint(2) | YES | | NA | |
| day | smallint(2) | YES | | NA | |
| dep_time | smallint(4) | YES | | NA | |
| sched_dep_time | smallint(4) | YES | | NA | |
| dep_delay | smallint(4) | YES | | NA | |
| arr_time | smallint(4) | YES | | NA | |
| sched_arr_time | smallint(4) | YES | | NA | |
| arr_delay | smallint(4) | YES | | NA | |
| carrier | varchar(2) | NO | | | |

This output is conceptually similar to R's `glimpse()` function, but it exposes the underlying database types (e.g., `varchar`, `int`, `datetime`). Databases require strict typing to optimize physical storage on disk. A column defined as `varchar(3)` can only ever hold text strings up to three characters long, preventing data corruption and saving immense amounts of storage space.

Remember, earlier we were trying to figure out if it would be more efficient to store student names versus their IDs!

To fully appreciate the power of lazy evaluation, we can compare the memory footprint of a database pointer versus a materialized in-memory object in R. Let us create a reference to the `carriers` table.

```
# reference to carriers table from the airlines database

carriers <- dplyr::tbl(con_air, "carriers")
```

```
DESCRIBE carriers;
```

Table 1.3: 2 records

| Field | Type | Null | Key | Default | Extra |
|---------|--------------|------|-----|---------|-------|
| carrier | varchar(7) | NO | PRI | | |
| name | varchar(255) | NO | | | |

Let's check how much RAM this connection occupies in our local environment:

```
# carriers lives in the SQL database and is linked remotely

carriers |>
  utils::object.size() |>
  print(units = "Kb")
```

7.4 Kb

The object size is remarkably small (often just a few kilobytes). It only stores the connection string and metadata, not the data itself.

However, in a typical data science workflow, we eventually need the data in our local environment—usually after we have aggregated or filtered it—so we can pass it to local functions like `ggplot2` for visualization. To bridge the gap between the remote database and our local R session, we use the `collect()` function.

```
carriers_df <- carriers |> dplyr::collect()
```

Let's check the memory size now:

```
# carriers lives in memory in R as carriers_df

carriers_df |>
  utils::object.size() |>
  print(units = "Kb")
```

234.8 Kb

The materialized data frame is significantly larger because the actual data now resides in our computer's RAM.

The Danger of `collect()`

We should never run `collect()` on a massive, unaggregated table like `flights`. That table contains hundreds of millions of rows. If we attempt to `collect()` the entire table, we will exhaust our machine's RAM, and our R session will crash. The proper workflow is to use `dbplyr` or `SQL` to filter and summarize the data on the remote server, and only `collect()` the finalized, summarized results into R. Another safer alternative (if needed) is

```
# flights_df lives in memory in R

flights_df <- flights |>
  head(5) |>
  collect()
```

1.4 terminate connection

Just as we would close a file after editing it, we must explicitly sever our connection to a remote database when our analysis is complete. Open connections consume resources on both our local machine and the remote server.

```
# terminate the SQL connection

DBI::dbDisconnect(con_air, shutdown = TRUE)
```

Terminating connections is an essential habit for responsible computing, particularly when sharing server resources in an industry or academic environment.

2 Writing SQL Queries

In our previous chapter, we established how to connect R to a remote relational database. However, simply connecting to a database is only the first step. Our primary goal is to extract, filter, and aggregate data to answer analytical questions.

Because we are operating within the R ecosystem but interacting with a **SQL** database, we have three distinct methods for sending instructions to the server:

1. Translating R's **dplyr** syntax into **SQL** using the **dbplyr** package.
2. Passing raw **SQL** strings through R functions using the **DBI** package.
3. Writing native **SQL** directly within Quarto code chunks.

In this chapter, we will explore all three methodologies, gradually transitioning from the familiar **tidyverse** syntax to writing robust, native **SQL**.

2.1 load packages and connect to database server

We begin by loading our required packages and establishing our connection to the **airlines** database, just as we did previously. We will also create a reference pointer to the **flights** table, which contains millions of commercial flight records.

```
library(tidyverse)
library(mdsr)
library(DBI)
library(RMariaDB)
library(dbplyr)
```

```
# connect to the airlines database

con_air <- mdsr::dbConnect_scidb("airlines")

DBI::dbListTables(con_air) # list tables in database
```

```
[1] "airports"      "planes"        "carriers"      "flights_summary"
[5] "flights"
```

```
# reference to flights table from the airlines database

flights <- dplyr::tbl(con_air, "flights")
```

```
DESCRIBE flights;
```

Table 2.1: Displaying records 1 - 10

| Field | Type | Null | Key | Default | Extra |
|----------------|-------------|------|-----|---------|-------|
| year | smallint(4) | YES | | NA | |
| month | smallint(2) | YES | | NA | |
| day | smallint(2) | YES | | NA | |
| dep_time | smallint(4) | YES | | NA | |
| sched_dep_time | smallint(4) | YES | | NA | |
| dep_delay | smallint(4) | YES | | NA | |
| arr_time | smallint(4) | YES | | NA | |
| sched_arr_time | smallint(4) | YES | | NA | |
| arr_delay | smallint(4) | YES | | NA | |
| carrier | varchar(2) | NO | | | |

2.2 SQL in R: Part 1 - the ‘dbplyr’ R package will directly translate ‘dplyr’ code into SQL

For data scientists already comfortable with the `tidyverse`, the `dbplyr` package is an extraordinary tool. It acts as a translation engine, taking our standard `dplyr` verbs (`filter`, `arrange`, `select`, `mutate`, `summarize`, `group_by`) and silently converting them into the specific SQL dialect required by our database (in this case, MariaDB/MySQL).

Suppose we want to find the minimum and maximum year recorded in our `flights` table. We can write standard `dplyr` code:

```
yrs <- flights |>
  summarize(min_year = min(year, na.rm = TRUE),
            max_year = max(year, na.rm = TRUE))

yrs
```

```
# A query: ?? x 2
# Database: mysql [mdsr_public@mdsr.crcbo51tmesf.us-east-2.rds.amazonaws.com:3306/airlines]
```

```

  min_year max_year
    <int>    <int>
1    2013    2015

```

The result `yrs` is a **lazy** query — it holds the SQL definition, not the data itself.

The data is processed on the remote server, and only the summarized one-row result is returned to our console. But how did R communicate this to the database? We can peek under the hood using the `show_query()` function. This reveals the SQL that `dbplyr` generated under the hood.

```
dbplyr::show_query(yrs)
```

```

<SQL>
SELECT MIN(`year`) AS `min_year`, MAX(`year`) AS `max_year`
FROM `flights`

```

The SQL Vocabulary

Looking at the output of `show_query()`, we see several foundational SQL keywords:

- **SELECT**: Determines which columns or calculations we want to retrieve.
- **AS**: Creates an alias, similar to the `=` sign in a `mutate()` or `summarize()` function.
- **FROM**: Specifies the table we are querying.

Notice also that the query ends with a semicolon (`;`).

Let us try a more complex example. We want to extract all flights from 2013 that traveled between Los Angeles (LAX) and Boston (BOS).

```

la_bos <- flights |>
  filter(year == 2013 &
         ((origin == "LAX" & dest == "BOS") |
          (origin == "BOS" & dest == "LAX")))

la_bos

```

```

# A query: ?? x 21
# Database: mysql [mdsr_public@mdsr.crcbo51tmesf.us-east-2.rds.amazonaws.com:3306/airlines]
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time

```

| | <int> | <int> | <int> | <int> | <int> | <int> | <int> | <int> |
|----|-------|-------|-------|-------|-------|-------|-------|-------|
| 1 | 2013 | 10 | 1 | 655 | 700 | -5 | 1508 | 1530 |
| 2 | 2013 | 10 | 1 | 704 | 710 | -6 | 1006 | 1035 |
| 3 | 2013 | 10 | 1 | 707 | 715 | -8 | 1000 | 1030 |
| 4 | 2013 | 10 | 1 | 755 | 800 | -5 | 1611 | 1626 |
| 5 | 2013 | 10 | 1 | 800 | 800 | 0 | 1054 | 1117 |
| 6 | 2013 | 10 | 1 | 804 | 800 | 4 | 1630 | 1630 |
| 7 | 2013 | 10 | 1 | 838 | 844 | -6 | 1706 | 1716 |
| 8 | 2013 | 10 | 1 | 1108 | 1110 | -2 | 1421 | 1425 |
| 9 | 2013 | 10 | 1 | 1202 | 1205 | -3 | 2015 | 2040 |
| 10 | 2013 | 10 | 1 | 1208 | 1209 | -1 | 2016 | 2035 |

```
# i more rows
# i 13 more variables: arr_delay <int>, carrier <chr>, tailnum <chr>,
#   flight <int>, origin <chr>, dest <chr>, air_time <int>, distance <int>,
#   cancelled <int>, diverted <int>, hour <int>, minute <int>, time_hour <dtm>
```

When we inspect the SQL translation for this filter, we notice some critical syntactical differences between R and SQL.

```
dplyr::show_query(la_bos)
```

```
<SQL>
SELECT *
FROM `flights`
WHERE (`year` = 2013.0 AND ((`origin` = 'LAX' AND `dest` = 'BOS') OR (`origin` = 'BOS' AND `dest` = 'LAX')))
```

In the SQL translation, our R `filter()` statement becomes a SQL `WHERE` clause. Furthermore, notice that SQL uses a single `=` for logical evaluation, whereas R requires `==`. Finally, our symbolic operators `&` and `|` are translated into explicit `AND` and `OR` keywords.

Once we are satisfied with our query and wish to bring the data into our local R environment for visualization or modeling, we append the `collect()` function. This should be done only when the result set is small enough to fit in memory.

```
la_bos <- la_bos |>
  collect()

la_bos
```

```
# A tibble: 7,227 x 21
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
```


	<int>	<int>	<int>	<int>	<int>	<int>	<int>	<int>
1	2013	10	1	655	700	-5	1508	1530
2	2013	10	1	704	710	-6	1006	1035
3	2013	10	1	707	715	-8	1000	1030
4	2013	10	1	755	800	-5	1611	1626
5	2013	10	1	800	800	0	1054	1117
6	2013	10	1	804	800	4	1630	1630
7	2013	10	1	838	844	-6	1706	1716
8	2013	10	1	1108	1110	-2	1421	1425
9	2013	10	1	1202	1205	-3	2015	2040
10	2013	10	1	1208	1209	-1	2016	2035

```
# i 7,217 more rows
# i 13 more variables: arr_delay <int>, carrier <chr>, tailnum <chr>,
#   flight <int>, origin <chr>, dest <chr>, air_time <int>, distance <int>,
#   cancelled <int>, diverted <int>, hour <int>, minute <int>, time_hour <dtm>
```

⚠ Where dbplyr Translation Breaks Down

Notice the code below. Depending on the database backend, this either fails outright or produces a warning, because `cumsum()` requires a **SQL window function**, and not every `dplyr` function has a clean, universal SQL equivalent — especially once we leave basic aggregation and move into ordered, row-by-row calculations. This is precisely the kind of situation where dropping down to raw SQL becomes necessary rather than a stylistic choice. In the sections that follow, we introduce two ways to do exactly that.

```
# attempting a window function (running total) using dplyr syntax
```

```
dbplyr_fail <- flights |>
  filter(origin == "ATW") |>
  group_by(year) |>
  summarize(num_flights = n()) |>
  mutate(running_total = cumsum(num_flights))
```

```
dbplyr::show_query(dbplyr_fail)
```

```
Warning: Windowed expression `SUM(`num_flights`)` does not have explicit order.
i Please use `arrange()`, `window_order()`, or `.order` to make deterministic.
```

```
<SQL>
```

```
SELECT *, SUM(`num_flights`) OVER (ROWS UNBOUNDED PRECEDING) AS `running_total`
FROM (
```

```

SELECT `year`, COUNT(*) AS `num_flights`
FROM `flights`
WHERE (`origin` = 'ATW')
GROUP BY `year`
) AS `q01`

```

2.3 SQL in R: Part 2 - using the DBI R package to write SQL queries in r code chunks

While dbplyr is incredibly convenient, it cannot translate every possible SQL function, especially highly specialized window functions or dialect-specific commands. In these instances, we can bypass the translation layer and pass raw SQL strings directly to the database using `DBI::dbGetQuery()`.

```

# retrieves the first 8 rows from flights

DBI::dbGetQuery(con_air,
  "SELECT * FROM flights LIMIT 8;")

```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
1	2013	10	1	2	10	-8	453	505
2	2013	10	1	4	2359	5	730	729
3	2013	10	1	11	15	-4	528	530
4	2013	10	1	14	2355	19	544	540
5	2013	10	1	16	17	-1	515	525
6	2013	10	1	22	20	2	552	554
7	2013	10	1	29	35	-6	808	816
8	2013	10	1	29	35	-6	449	458

	arr_delay	carrier	tailnum	flight	origin	dest	air_time	distance	cancelled
1	-12	AA	N201AA	2400	LAX	DFW	149	1235	0
2	1	FL	N344AT	710	SFO	ATL	247	2139	0
3	-2	AA	N3KMAA	1052	SFO	DFW	182	1464	0
4	4	AA	N3ENAA	2392	SEA	ORD	191	1721	0
5	-10	UA	N38473	1614	LAX	IAH	157	1379	0
6	-2	UA	N458UA	291	SFO	IAH	188	1635	0
7	-8	US	N551UW	436	LAX	CLT	256	2125	0
8	-9	AS	N402AS	108	ANC	SEA	181	1448	0

	diverted	hour	minute	time_hour
1	0	0	10	2013-10-01 00:10:00

2	0	23	59	2013-10-01	23:59:00
3	0	0	15	2013-10-01	00:15:00
4	0	23	55	2013-10-01	23:55:00
5	0	0	17	2013-10-01	00:17:00
6	0	0	20	2013-10-01	00:20:00
7	0	0	35	2013-10-01	00:35:00
8	0	0	35	2013-10-01	00:35:00

In SQL, the `*` symbol is a wildcard meaning “all columns.” The `LIMIT` clause is the SQL equivalent of R’s `head()` function, restricting the output size.

We can also perform complex aggregations directly within the string. Here, we count the number of flights per year and order the results.

```
DBI::dbGetQuery(con_air,
  "SELECT year, count(*) AS num_flights FROM flights GROUP BY year ORDER BY num_flights;")
```

	year	num_flights
1	2015	5819079
2	2014	5819811
3	2013	6369482

2.4 SQL in R: Part 3 - writing SQL queries in sql code chunks

Passing SQL strings inside R functions can become visually cumbersome for long, multi-line queries. Because Quarto is a polyglot environment, we can change the code chunk type from `{r}` to `{sql}`. This provides proper syntax highlighting and allows us to write native SQL exactly as we would in a dedicated database client.

Let us recreate our previous `dbplyr` queries using native SQL chunks.

```
-- display minimum and maximum year

SELECT MIN(year) AS min_year,
       MAX(year) AS max_year
FROM flights;
```

Table 2.2: 1 records

min_year	max_year
2013	2015

```
-- extract flights between LAX and BOS in 2013

SELECT *
FROM flights
WHERE (year = 2013 AND
      ((origin = 'LAX' AND dest = 'BOS') OR
       (origin = 'BOS' AND dest = 'LAX')));

-- notice the use of = (not ==) for equality comparisons
```

Table 2.3: Displaying records 1 - 10

year	month	day	dep_delay	dep_time	dep_delay_min	dep_delay_max	arr_delay	arr_time	tailnum	flight	origin	dest	aircraft	distance	cancelled	diverted	hour	minute	time_hour
2013	0	1	655	700	-5	1508	1530	-	AA N5DP	22	LAX	BOS	286	2611	0	0	7	0	2013-10-01 07:00:00
2013	0	1	704	710	-6	1006	1035	-	AA N3G2	29	BOS	LAX	337	2611	0	0	7	10	2013-10-01 07:10:00
2013	0	1	707	715	-8	1000	1030	-	VX N837	30	BOS	LAX	329	2611	0	0	7	15	2013-10-01 07:15:00
2013	0	1	755	800	-5	1611	1626	-	B6 N605	15	LAX	BOS	300	2611	0	0	8	0	2013-10-01 08:00:00
2013	0	1	800	800	0	1054	1117	-	B6 N612	23	BOS	LAX	331	2611	0	0	8	0	2013-10-01 08:00:00

year	month	day	dep_scheduled	deptime	dep_time	arr_scheduled	arr_time	carrier	tailnum	origin	flight	in	dest	air_time	cancelled	diverted	hours	minutes	time_hour
2013	10	1	804	800	4	1630	1630	0	VX N633CA	LAX	BOS	10	2611	0	0	8	0	2013-10-01 08:00:00	
2013	10	1	838	844	-6	1706	1716 10	-	UA N515UA	LAX	BOS	299	2611	0	0	8	44	2013-10-01 08:44:00	
2013	10	1	1108	1110	-2	1421	1425	-4	VX N633CA	BOS	LAX	347	2611	0	0	11	10	2013-10-01 11:10:00	
2013	10	1	1202	1205	-3	2015	2040 25	-	AA N3JU	LAX	BOS	291	2611	0	0	12	5	2013-10-01 12:05:00	
2013	10	1	1208	1209	-1	2016	2035 19	-	B6 N612JB	LAX	BOS	291	2611	0	0	12	9	2013-10-01 12:09:00	

LIMIT restricts the number of rows returned. This is useful for previewing data without pulling the entire table from the server.

```
-- select all columns and limit to top 2 rows

SELECT *
FROM flights
LIMIT 2;
```

Table 2.4: 2 records

year	month	day	dep_scheduled	deptime	dep_time	arr_scheduled	arr_time	carrier	tailnum	origin	flight	in	dest	air_time	cancelled	diverted	hours	minutes	time_hour
2013	10	1	2	10	-8	453	505	-	AA N20240	LAX	DFW	49	1235	0	0	0	10	2013-10-01 00:10:00	

year	month	day	dep_scheduled	dep_time	dep_delay	dep_delay_min	arr_time	arr_delay	tailnum	flight	origin	destination	air_time	cancelled	diverted	count	hour
2013	10	1	4	2359	5	730	729	1	FL N347AT	410	SFO	ATL	247	2139	0	0	23 59
																	2013-10-01 23:59:00

GROUP BY partitions rows into groups sharing the same value of one or more columns, and aggregate functions like **COUNT()** are applied within each group. **ORDER BY** sorts the result; the default sort order is ascending (use **ORDER BY num_flights DESC** for descending order).

```
-- select year, and the computing num_flights from the flights table
-- grouping by year
-- ordering by the num_flights
```

```
SELECT year,
        COUNT(*) AS num_flights
FROM flights
GROUP BY year
ORDER BY num_flights;
```

Table 2.5: 3 records

year	num_flights
2015	5819079
2014	5819811
2013	6369482

Written Order of SQL Clauses

Queries in SQL start with the **SELECT** keyword and consist of several **clauses**, which have to be written in this order:

- **SELECT** allows you to list the columns, or functions operating on columns, that you want to retrieve.
- **FROM** specifies the table where the data are.
- **JOIN** allows you to stitch together two or more tables using a key.
- **WHERE** allows you to extract rows according to some criteria.

- **GROUP BY** allows you to aggregate the records according to some shared value.
- **HAVING** is like a **WHERE** clause that operates on the result set—not the original rows themselves.
- **ORDER BY** orders the rows of the result set.
- **LIMIT** restricts the number of rows in the output.

The following queries explore the `flights` table, starting with basic retrieval and building toward more complex aggregations.

```
DESCRIBE flights;
```

Table 2.6: Displaying records 1 - 10

Field	Type	Null	Key	Default	Extra
year	smallint(4)	YES		NA	
month	smallint(2)	YES		NA	
day	smallint(2)	YES		NA	
dep_time	smallint(4)	YES		NA	
sched_dep_time	smallint(4)	YES		NA	
dep_delay	smallint(4)	YES		NA	
arr_time	smallint(4)	YES		NA	
sched_arr_time	smallint(4)	YES		NA	
arr_delay	smallint(4)	YES		NA	
carrier	varchar(2)	NO			

2.4.1 query 1

`SELECT *` retrieves all columns. Combined with `LIMIT`, this is a quick way to preview the contents of a table without loading all rows.

```
-- display the first ten rows of the flights table

SELECT *
FROM flights
LIMIT 10;
```

Table 2.7: Displaying records 1 - 10

year	month	day	dep_scheduled	dep_actual	dep_time	arr_scheduled	arr_time	carrier	tailnum	flight	origin	destination	airline	distance	cancelled	diverted	hours	minutes	time_hour
2013	0	1	2	10	-8	453	505	-	AA	N2024	LAX	DFW	49	1235	0	0	0	10	2013-10-01 00:10:00
2013	0	1	4	2359	5	730	729	1	FL	N347AT	SFO	ATL	247	2139	0	0	23	59	2013-10-01 23:59:00
2013	0	1	11	15	-4	528	530	-2	AA	N3K13	MIA	FODFW	82	1464	0	0	0	15	2013-10-01 00:15:00
2013	0	1	14	2355	19	544	540	4	AA	N3E32	SEA	ORD	91	1721	0	0	23	55	2013-10-01 23:55:00
2013	0	1	16	17	-1	515	525	-	UA	N3847	LAX	IAH	57	1379	0	0	0	17	2013-10-01 00:17:00
2013	0	1	22	20	2	552	554	-2	UA	N4589	AS	FOIAH	88	1635	0	0	0	20	2013-10-01 00:20:00
2013	0	1	29	35	-6	808	816	-8	US	N5543	SVLAX	CLT	256	2125	0	0	0	35	2013-10-01 00:35:00
2013	0	1	29	35	-6	449	458	-9	AS	N4021	AS	ANSEA	81	1448	0	0	0	35	2013-10-01 00:35:00
2013	0	1	31	30	1	519	538	-	DL	N5968	MS	SEAMS	47	1399	0	0	0	30	2013-10-01 00:30:00

year	month	day	dep_scheduled	dep_time	dep_time	arr_time	carrier	tailnum	flight	origin	destination	air_time	cancelled	diverted	hours	minutes	time_hour	
2013	10	1	32	33	-1	557	606	-9	DL N14081	LAX	MSP	83	1535	0	0	0	33	2013-10-01 00:33:00

Essentially, LIMIT takes two arguments - the first skips a specific number of rows (offset), the second returns the specified row count after skipping.

```
-- skip the first 18,008,362 rows and return the next 10
-- how did we find out the total number of rows?

SELECT *
FROM flights
LIMIT 18008362, 10;
```

Table 2.8: Displaying records 1 - 10

year	month	day	dep_scheduled	dep_time	dep_time	arr_time	carrier	tailnum	flight	origin	destination	air_time	cancelled	diverted	hours	minutes	time_hour	
2015	9	30	0	930	0	0	1100	0	WN N65308	PHX	LAX	0	370	1	0	9	30	2015-09-30 09:30:00
2015	9	30	0	1100	0	0	1340	0	WN N40205	SEA	SNA	0	978	1	0	11	0	2015-09-30 11:00:00
2015	9	30	0	735	0	0	1025	0	WN N40205	SEA	SEA	0	978	1	0	7	35	2015-09-30 07:35:00
2015	9	30	0	600	0	0	850	0	NK N60521	KORD	ATL	0	606	1	0	6	0	2015-09-30 06:00:00

year	month	day	dep_scheduled	dep_delay	dep_time	arr_scheduled	arr_delay	carrier	tailnum	flight	origin	destination	aircraft	distance	cancelled	diverted	hours	minutes	time_hour
2015	9	30	0	940	0	0	1050	0	NK	N6053K	ATL	IAH	0	689	1	0	9	40	2015-09-30 09:40:00
2015	9	30	0	1955	0	0	2256	0	NK	N640NN	LGA	FLL	0	1076	1	0	19	55	2015-09-30 19:55:00
2015	9	30	0	1054	0	0	1228	0	EV	N9162Z	AOA	JATL	0	399	1	0	10	54	2015-09-30 10:54:00
2015	9	30	0	1950	0	0	2035	0	OO	N8752L	PHX	FLG	0	119	1	0	19	50	2015-09-30 19:50:00
2015	9	30	0	620	0	0	804	0	OO	N8752L	SBP	PHX	0	509	1	0	6	20	2015-09-30 06:20:00
2015	9	30	0	1145	0	0	2025	0	UA	NA 1766	SFO	WR	0	2565	1	0	11	45	2015-09-30 11:45:00

2.4.2 query 2

Rather than retrieving all columns, we can name specific columns in the **SELECT** clause. This reduces the amount of data transferred from the server and makes results easier to interpret.

```
-- display year, month, day, origin, dep_time, sched_dep_time, dep_delay for the first ten rows
SELECT year, month, day, origin, dep_time, sched_dep_time, dep_delay
FROM flights
LIMIT 10;
```

Table 2.9: Displaying records 1 - 10

year	month	day	origin	dep_time	sched_dep_time	dep_delay
2013	10	1	LAX	2	10	-8
2013	10	1	SFO	4	2359	5
2013	10	1	SFO	11	15	-4
2013	10	1	SEA	14	2355	19
2013	10	1	LAX	16	17	-1
2013	10	1	SFO	22	20	2
2013	10	1	LAX	29	35	-6
2013	10	1	ANC	29	35	-6
2013	10	1	SEA	31	30	1
2013	10	1	LAX	32	33	-1

We can also concatenate strings and cast them into date formats natively.

```
SELECT STR_TO_DATE(CONCAT(year, '-', month, '-', day), '%Y-%m-%d') AS 'date', origin, dep_time
FROM flights
LIMIT 10;
```

Table 2.10: Displaying records 1 - 10

date	origin	dep_time	sched_dep_time	dep_delay
2013-10-01	LAX	2	10	-8
2013-10-01	SFO	4	2359	5
2013-10-01	SFO	11	15	-4
2013-10-01	SEA	14	2355	19
2013-10-01	LAX	16	17	-1
2013-10-01	SFO	22	20	2
2013-10-01	LAX	29	35	-6
2013-10-01	ANC	29	35	-6
2013-10-01	SEA	31	30	1
2013-10-01	LAX	32	33	-1

2.4.3 query 3

COUNT(*) counts every row in the table regardless of NULL values in individual columns. Aliasing the result with AS gives the output column a descriptive name.

```
-- count the total number of rows in the flights table

SELECT COUNT(*) AS total_rows
FROM flights;
```

Table 2.11: 1 records

total_rows
18008372

COUNT(*) counts *rows*, but we often want to count **unique values** instead — for example, how many different carriers are in this table. In R, `n_distinct()` does this; in SQL, we pair COUNT() with the DISTINCT keyword.

```
SELECT COUNT(DISTINCT carrier) AS num_carriers
FROM flights;
```

Table 2.12: 1 records

num_carriers
17

2.4.4 query 4

WHERE filters rows before any aggregation takes place. Here we restrict to flights from ATW to ORD in 2013, combining multiple conditions with AND.

```
-- extract all info about flights that flew from Appleton (ATW) to Chicago (ORD) in 2013

SELECT *
FROM flights
WHERE origin = 'ATW' AND dest = 'ORD' AND year = 2013;
```

Table 2.13: Displaying records 1 - 10

year	month	day	dep_scheduled	dep_actual	dep_delay	deptime	arr_scheduled	arr_actual	arr_delay	carrier	tailnum	flight	origin	destination	airline	distance	cancelled	diverted	hours	minutes	time_hour
2013	10	1	543	545	-2	633	640	-7	EV	N1438	9082	ATW	ORB5	160	0	0	5	45	2013-10-01	05:45:00	
2013	10	1	927	928	-1	1027	1023	4	OO	N9175	550	ATW	ORB7	160	0	0	9	28	2013-10-01	09:28:00	
2013	10	1	1137	1137	0	1240	1230	10	EV	N1438	9082	ATW	ORB9	160	0	0	11	37	2013-10-01	11:37:00	
2013	10	1	1312	1322	-10	1404	1417	-13	EV	N1551	551	ATW	ORB6	160	0	0	13	22	2013-10-01	13:22:00	
2013	10	1	1619	1627	-8	1728	1724	4	EV	N1299	927	ATW	ORB3	160	0	0	16	27	2013-10-01	16:27:00	
2013	10	1	1818	1823	-5	1920	1922	-2	EV	N1551	551	ATW	ORB3	160	0	0	18	23	2013-10-01	18:23:00	
2013	10	2	542	545	-3	631	640	-9	EV	N1438	9082	ATW	ORB3	160	0	0	5	45	2013-10-02	05:45:00	
2013	10	2	923	928	-5	1018	1023	-5	OO	N9205	550	ATW	ORB7	160	0	0	9	28	2013-10-02	09:28:00	
2013	10	2	1127	1137	-10	1212	1230	-18	EV	N1358	582	ATW	ORB3	160	0	0	11	37	2013-10-02	11:37:00	

year	month	day	dep_scheduled	dep_time	dep_delay	dep_delay_min	arr_scheduled	arr_time	arr_delay	tailnum	flight	origin	dest	air_carrier	time_cancelled	diverted	hours	minutes	time_hour	
2013	10	2	130713	2215	-	1353	1417	-	EV	N13917	917	ATW	ORD	8	160	0	0	13	22	2013-10-02 13:22:00

i Logical Operators AND and OR

Logical operators in SQL evaluate in a strict order: **AND** is evaluated before **OR**. Imagine we are looking at our **flights** table and we want to find all flights in the year 2013 that either departed from Appleton (ATW) or arrived in Chicago (ORD). It is tempting to write the query exactly as we speak it:

```
SELECT COUNT(*) AS num_flights
FROM flights
WHERE year = 2013 AND origin = 'ATW' OR dest = 'ORD';
```

Table 2.14: 1 records

<u>num_flights</u>
910681

Because AND has higher precedence, the SQL engine secretly groups the query like this:

```
SELECT COUNT(*)  
FROM flights  
WHERE (year = 2013 AND origin = 'ATW') OR dest = 'ORD';
```

Table 2.15: 1 records

<u>COUNT(*)</u>
910681

What this actually returns:

1. Every flight that departed from ATW in 2013.
2. PLUS every single flight that landed in ORD, regardless of the year, and regardless of where it came from.

Instead of returning a small subset of specific 2013 flights, our query returns millions of rows, pulling in ORD flights from 2001, 2010, etc.

To fix this, we use parentheses to force the database to evaluate the **OR** condition as a single, combined logical unit before it evaluates the **AND** condition.

```
SELECT COUNT(*) AS num_flights
FROM flights
WHERE year = 2013 AND (origin = 'ATW' OR dest = 'ORD');
```

Table 2.16: 1 records

num_flights
309790

How SQL evaluates this:

1. First, it looks inside the parentheses: “Is the origin ATW, OR is the destination ORD?” (Returns True or False for each row).
2. Then, it evaluates the **AND**: “Is the year 2013 AND did the previous step return True?”

By wrapping the **OR** statements in brackets, we ensure that the **year = 2013** condition acts as a strict filter applied to both the origin and the destination checks.

A Golden Rule for SQL:

Whenever you have a WHERE clause that mixes AND and OR, always use parentheses. Even if you memorize the precedence rules, parentheses make your code instantly readable to your colleagues and future self, eliminating any ambiguity about your analytical intentions.

BETWEEN vs IN

Notice the difference between the **BETWEEN** and **IN** operators.


```
-- number of flights that were delayed at departure by less than 10 minutes

SELECT COUNT(*) as num_rows
FROM flights
WHERE dep_delay BETWEEN 0 AND 10;
```

Table 2.17: 1 records

num_rows
4140140

```
-- number of flights that were delayed by either 0 or 10 minutes

SELECT COUNT(*) as num_rows
FROM flights
WHERE dep_delay IN (0,10);
```

Table 2.18: 1 records

num_rows
1471502

2.4.5 query 5

Adding `GROUP BY year` counts the number of flights that operated on the ATW–ORD route per year.

`ORDER BY year` sorts the results chronologically.

```
-- count the number of flights from Appleton (ATW) for each year

SELECT year, COUNT(*) AS num_flights
FROM flights
WHERE origin = 'ATW' AND dest = 'ORD'
GROUP BY year
ORDER BY year;
```

Table 2.19: 3 records

year	num_flights
2013	1838
2014	1647
2015	1570

We can group by multiple variables by separating them with commas. Grouping by both **carrier** and **year** lets us see how many flights each airline operated on the ATW–ORD route per year.

```
-- which carriers flew from Appleton (ATW) to Chicago (ORD)
-- how many flights did they do per year

SELECT carrier, year, COUNT(*) AS num_flights
FROM flights
WHERE origin = 'ATW' AND dest = 'ORD'
GROUP BY carrier, year
ORDER BY carrier, year;
```

Table 2.20: 6 records

carrier	year	num_flights
EV	2013	1634
EV	2014	1135
EV	2015	1297
OO	2013	204
OO	2014	512
OO	2015	273

2.4.6 query 6

MAX(dep_time) within each carrier group returns the latest recorded departure time for that carrier on the ATW–ORD route.

Note that departure times in this table are stored as integers in HHMM format (e.g., 1435 = 2:35 PM).

```
-- for flights that flew from Appleton (ATW) to Chicago (ORD)
-- what was the latest flight for each carrier

SELECT carrier, MAX(dep_time) AS latest_dep_time
FROM flights
WHERE origin = 'ATW' AND dest = 'ORD'
GROUP BY carrier
LIMIT 1;
```

Table 2.21: 1 records

carrier	latest_dep_time
EV	2309

2.4.7 query 7

Understanding the difference between **WHERE** and **HAVING** is one of the most important hurdles in learning **SQL**. Both act as filters, but they operate at different stages of the execution pipeline.

WHERE filters records in the original dataset before any grouping occurs. **HAVING** filters the results of our aggregations after the **GROUP BY** clause has executed.

First, let us find the destination with the lowest average arrival delay from Appleton, using **WHERE** to filter the origin.

```
-- for flights that flew from Appleton (ATW)
-- which destination had the lowest average arrival delay time

SELECT dest, COUNT(*) AS num_flights, AVG(arr_delay) AS avg_arr_delay
FROM flights
WHERE origin = 'ATW'
GROUP BY dest
ORDER BY avg_arr_delay;
```

Table 2.22: 4 records

dest	num_flights	avg_arr_delay
DTW	2183	-0.6949
MSP	1920	0.9250
ATL	1756	6.5017

dest	num_flights	avg_arr_delay
ORD	5055	13.9937

Suppose we want to restrict our analysis to robust routes — say, those with at least 2,000 flights. We cannot put `num_flights >= 2000` in the `WHERE` clause because `num_flights` does not exist in the original table; it is an aggregated calculation. Therefore, we must use `HAVING`.

```
-- for flights that flew from Appleton (ATW)
-- consider only destinations with at least 2000 flights from Appleton (ATW)
-- which destination had the lowest average arrival delay time

SELECT dest, COUNT(*) AS num_flights, AVG(arr_delay) AS avg_arr_delay
FROM flights
WHERE origin = 'ATW'
GROUP BY dest
HAVING num_flights >= 2000
ORDER BY avg_arr_delay;
```

Table 2.23: 2 records

dest	num_flights	avg_arr_delay
DTW	2183	-0.6949
ORD	5055	13.9937

`HAVING` filters groups **after** aggregation, unlike `WHERE` which filters rows **before** aggregation.

2.4.8 query 8

We now have all the pieces to write a complete query using every clause introduced in this chapter. Notice how each clause does a distinct job, and how the order we write them in matches the order listed at the start of this section: we filter raw rows first (`WHERE`), then aggregate (`GROUP BY`), then filter the aggregated results (`HAVING`), then sort (`ORDER BY`), and finally restrict the output size (`LIMIT`). Try reading the query out loud, clause by clause — it should read almost like a sentence describing the question we are asking of the data.

```
-- for flights out of Appleton (ATW) in 2013 or later
-- consider only destinations with at least 500 flights
-- find the 5 destinations with the highest average departure delay
```

```

SELECT dest,
       COUNT(*) AS num_flights,
       AVG(dep_delay) AS avg_dep_delay
FROM flights
WHERE origin = 'ATW' AND year >= 2013
GROUP BY dest
HAVING num_flights >= 500
ORDER BY avg_dep_delay DESC
LIMIT 5;

```

Table 2.24: 4 records

dest	num_flights	avg_dep_delay
ORD	5055	11.7658
ATL	1756	9.4129
DTW	2183	3.5456
MSP	1920	2.2849

i Which Approach to Use?

Having now seen all three approaches to writing SQL from within R, a natural question is: which one should we actually use? There is no single right answer, but the following rough guide is a useful starting point:

- **Use `dbplyr` (Part 1)** when your query is a fairly standard `dplyr` pipeline — filtering, grouping, summarizing — and you want your code to stay readable and consistent with the rest of a `tidyverse`-based analysis. This is the best default for most exploratory work.
- **Use `DBI::dbGetQuery()` (Part 2)** when you need a SQL feature `dbplyr` cannot translate (such as window functions), but the query is short enough to comfortably live as a string inside an R script or function.
- **Use native `{sql}` code chunks (Part 3)** when writing long, complex, or highly SQL-specific queries — window functions, subqueries, or multi-table joins — where syntax highlighting and formatting readability matter, or when you want a query that could be copy-pasted directly into a standalone database client.

In practice, many real analyses mix all three approaches within the same project, choosing whichever tool best fits the query at hand.

2.5 calling SQL generated datasets into R for data viz

Our ultimate workflow often involves doing the heavy lifting (filtering and aggregating) in SQL, and then pulling the summarized dataset into R to leverage packages like `ggplot2`. We have three ways to save the output of a SQL query into an R object:

2.5.1 Method 1: using `DBI::dbGetQuery()`

```
df1 <- DBI::dbGetQuery(con_air, "SELECT * FROM flights LIMIT 10")

# always check the number of rows
```

2.5.2 Method 2: using `dbplyr` with raw SQL

```
df2 <- dplyr::tbl(con_air, dplyr::sql("SELECT * FROM flights LIMIT 10")) |> dplyr::collect()

# doesn't need a semicolon
```

2.5.3 Method 3: Quarto chunk options

We can write a native SQL chunk and use the `--| output.var: "df"` code chunk option to assign the result to an object in our R environment.

```
SELECT * FROM flights LIMIT 10;
```

All three methods are valid. We select the one that best fits our workflow and the complexity of the query we need to execute.

2.6 terminate connection

Finally, always remember to disconnect your connection to the server.

```
# terminate the SQL connection

DBI::dbDisconnect(con_air, shutdown = TRUE)
```

3 Joins, Unions, and Subqueries

In our previous chapters, we focused on extracting and aggregating data from a single table. However, the defining feature of a **relational database** is that data is deliberately distributed across multiple tables. This design principle, known as **normalization**, eliminates data redundancy, saves storage space, and ensures data integrity.

For example, instead of repeating the full name of an airline on every single row of the `flights` dataset, the `flights` table only stores a brief carrier code (e.g., “DL”). If we want the full name (“Delta Air Lines Inc.”), we must link the `flights` table to the `carriers` table.

In this chapter, we explore how to stitch these normalized tables back together.

3.1 load packages and connect to database server

Let us load our required packages and establish our connection to the `airlines` database.

```
library(tidyverse)
library(mdsr)
library(DBI)
library(RMariaDB)
library(dbplyr)
```

```
# access airlines database
```

```
con_air <- mdsr::dbConnect_scldb("airlines")
```

```
DBI::dbListTables(con_air) # list tables in database
```

```
[1] "airports"      "planes"        "carriers"      "flights_summary"
[5] "flights"
```

We will create our standard lazy pointers for the three tables we will analyze in this chapter: `flights`, `carriers`, and `airports`.

```
# access tables from the airlines database

flights <- dplyr::tbl(con_air, "flights")

carriers <- dplyr::tbl(con_air, "carriers")

airports <- dplyr::tbl(con_air, "airports")
```

Before joining tables, it is standard practice to examine their schemas using the `DESCRIBE` command. We must identify the **keys** — the columns that exist in both tables and will serve as our connective tissue.

3.1.1 flights table

```
DESCRIBE flights;
```

Table 3.1: Displaying records 1 - 10

Field	Type	Null	Key	Default	Extra
year	smallint(4)	YES		NA	
month	smallint(2)	YES		NA	
day	smallint(2)	YES		NA	
dep_time	smallint(4)	YES		NA	
sched_dep_time	smallint(4)	YES		NA	
dep_delay	smallint(4)	YES		NA	
arr_time	smallint(4)	YES		NA	
sched_arr_time	smallint(4)	YES		NA	
arr_delay	smallint(4)	YES		NA	
carrier	varchar(2)	NO			

```
SELECT * FROM flights LIMIT 10;
```


Table 3.2: Displaying records 1 - 10

year	month	day	dep_sch	dep_time	dep_lat	dep_lon	arr_time	carrier	tailnum	flight	origin	destination	airline	distance	cancelled	diverted	hours	minutes	time_hour
2013	0	1	2	10	-8	453	505	-	AA	N2024	LAX	DFW	49	1235	0	0	0	10	2013-10-01 00:10:00
2013	0	1	4	2359	5	730	729	1	FL	N347AT	SFO	ATL	247	2139	0	0	23	59	2013-10-01 23:59:00
2013	0	1	11	15	-4	528	530	-2	AA	N3KM	MIA	FODFW	82	1464	0	0	0	15	2013-10-01 00:15:00
2013	0	1	14	2355	19	544	540	4	AA	N3E239	SEA	ORD	91	1721	0	0	23	55	2013-10-01 23:55:00
2013	0	1	16	17	-1	515	525	-	UA	N3847	LAX	IAH	57	1379	0	0	0	17	2013-10-01 00:17:00
2013	0	1	22	20	2	552	554	-2	UA	N4589	ASFO	IAH	88	1635	0	0	0	20	2013-10-01 00:20:00
2013	0	1	29	35	-6	808	816	-8	US	N5543	SVLAX	CLT	256	2125	0	0	0	35	2013-10-01 00:35:00
2013	0	1	29	35	-6	449	458	-9	AS	N402AS	ANCSEA	SEA	81	1448	0	0	0	35	2013-10-01 00:35:00
2013	0	1	31	30	1	519	538	-	DL	N5968	MSSEA	MSP	47	1399	0	0	0	30	2013-10-01 00:30:00

year	month	day	dep_scheduled	dep_time	dep_time	arr_time	carrier	tail	origin	destination	air_time	cancelled	diverted	hours	minutes	time_hour			
2013	10	1	32	33	-1	557	606	-9	DL	N1402A	LAX	MSP	83	1535	0	0	0	33	2013-10-01 00:33:00

3.1.2 carriers table

```
DESCRIBE carriers;
```

Table 3.3: 2 records

Field	Type	Null	Key	Default	Extra
carrier	varchar(7)	NO	PRI		
name	varchar(255)	NO			

```
SELECT * FROM carriers LIMIT 100;
```

Table 3.4: Displaying records 1 - 10

carrier	name
02Q	Titan Airways
04Q	Tradewind Aviation
05Q	Comlux Aviation, AG
06Q	Master Top Linhas Aereas Ltd.
07Q	Flair Airlines Ltd.
09Q	Swift Air, LLC
0BQ	DCA
0CQ	ACM AIR CHARTER GmbH
0GQ	Inter Island Airways, d/b/a Inter Island Air
0HQ	Polar Airlines de Mexico d/b/a Nova Air

3.1.3 airports table

```
DESCRIBE airports;
```

Table 3.5: 9 records

Field	Type	Null	Key	Default	Extra
faa	varchar(3)	NO	PRI		
name	varchar(255)	YES		NA	
lat	decimal(10,7)	YES		NA	
lon	decimal(10,7)	YES		NA	
alt	int(11)	YES		NA	
tz	smallint(4)	YES		NA	
dst	char(1)	YES		NA	
city	varchar(255)	YES		NA	
country	varchar(255)	YES		NA	

```
SELECT * FROM airports LIMIT 100;
```

Table 3.6: Displaying records 1 - 10

faa	name	lat	lon	alt	tz	dst	city	country
04G	Lansdowne Airport	41.13047	-80.61958	1044	-5	A	Youngstown	United States
06A	Moton Field Municipal Airport	32.46057	-85.68003	264	-6	A	Tuskegee	United States
06C	Schaumburg Regional	41.98934	-88.10124	801	-6	A	Schaumburg	United States
06N	Randall Airport	41.43191	-74.39156	523	-5	A	Middletown	United States
09J	Jekyll Island Airport	31.07447	-81.42778	11	-5	A	Jekyll Island	United States
0A9	Elizabethton Municipal Airport	36.37122	-82.17342	1593	-5	A	Elizabethton	United States
0G6	Williams County Airport	41.46731	-84.50678	730	-5	A	Bryan	United States
0G7	Finger Lakes Regional Airport	42.88356	-76.78123	492	-5	A	Seneca Falls	United States
0P2	Shoestring Aviation Airfield	39.79482	-76.64719	1000	-5	U	Stewartstown	United States

faa	name	lat	lon	alt	tz	dst	city	country
0S9	Jefferson County Intl	48.05381	- 122.81064	108	-8	A	Port Townsend	United States

3.2 Joins (Horizontal Integration)

To combine tables horizontally (adding columns from one table to another), we use **SQL** joins. In general, we must specify four pieces of information to successfully join two tables:

- name of the first table
- type of join
- name of the second table
- condition(s) of the join

The MySQL dialect supports several types of joins, which translate directly to the verbs we use in `dplyr`:

- **JOIN (or INNER JOIN)**: Returns only the rows that have matching values in both tables. Equivalent to `dplyr::inner_join()` in R.
- **LEFT JOIN** - Returns all rows from the left table, and the matched rows from the right table. Unmatched rows yield `NULL`. Equivalent to `dplyr::left_join()`.
- **RIGHT JOIN** - Returns all rows from the right table, and the matched rows from the left. Equivalent to `dplyr::right_join()`.
- **CROSS JOIN** - Returns the Cartesian product of rows from both tables. (Rarely used in standard data science workflows as it generates massive output).

The Missing FULL JOIN

Notice that `FULL JOIN` (equivalent to R's `full_join()`) is missing from this list. **MySQL** and **MariaDB** do not natively support full outer joins. If we were using **PostgreSQL** or **DuckDB**, `FULL OUTER JOIN` would be available. In **MySQL**, we have to construct a full join manually by taking a `LEFT JOIN` and combining it with a `RIGHT JOIN` using a `UNION`.

3.2.1 INNER JOIN

Suppose we want to access the full airport names for all flights that flew out of Appleton (ATW) in 2013. The `flights` table contains the destination code (`dest`), and the `airports` table contains the destination code (`faa`) alongside the full name.

```
-- access airport names for flights that flew out of Appleton in 2013

SELECT origin, dest, a.name AS dest_name, flight, carrier
FROM flights AS f
JOIN airports AS a ON f.dest = a.faa
WHERE year = 2013 AND origin = 'ATW'
LIMIT 10;
```

Table 3.7: Displaying records 1 - 10

origin	dest	dest_name	flight	carrier
ATW	ORD	Chicago Ohare Intl	6082	EV
ATW	ATL	Hartsfield Jackson Atlanta Intl	5476	EV
ATW	DTW	Detroit Metro Wayne Co	4078	9E
ATW	MSP	Minneapolis St Paul Intl	3870	9E
ATW	ORD	Chicago Ohare Intl	5650	OO
ATW	DTW	Detroit Metro Wayne Co	3694	9E
ATW	ORD	Chicago Ohare Intl	5852	EV
ATW	MSP	Minneapolis St Paul Intl	4593	OO
ATW	ORD	Chicago Ohare Intl	6131	EV
ATW	MSP	Minneapolis St Paul Intl	3798	9E

Notice our use of the **AS** keyword in the **FROM** and **JOIN** clauses to create aliases (**f** for **flights** and **a** for **airports**). This is a critical best practice. It saves us from typing the full table name repeatedly, and more importantly, it prevents ambiguity errors. If both tables contained a column named **year**, calling **SELECT year** would cause the query to crash; we would have to explicitly declare **SELECT f.year**.

Data architectures often require us to chain **multiple joins** together to assemble a complete picture. What if we want the full names of the origin airport, the destination airport, and the carrier?

Because we need airport names for both the origin and the destination, we must join the **airports** table to our query twice. To do this, we must give the **airports** table two distinct aliases.

```
-- access both airport and carrier names for flights that flew out of Appleton in 2013

SELECT origin, a2.name AS origin_name,
       dest, a1.name AS dest_name,
       flight,
```

```

        f.carrier, c.name AS carrier_name
FROM flights AS f
JOIN airports AS a1 ON f.dest = a1.faa
JOIN airports AS a2 ON f.origin = a2.faa
JOIN carriers AS c ON f.carrier = c.carrier
WHERE origin = 'ATW' AND year = 2013
LIMIT 10;

```

Table 3.8: Displaying records 1 - 10

origin	ori- gin_name	dest	dest_name	flight	carrier	carrier_name
ATW	Appleton	ORD	Chicago Ohare Intl	6082	EV	ExpressJet Airlines Inc.
ATW	Appleton	ATL	Hartsfield Jackson Atlanta Intl	5476	EV	ExpressJet Airlines Inc.
ATW	Appleton	DTW	Detroit Metro Wayne Co	4078	9E	Endeavor Air Inc.
ATW	Appleton	MSP	Minneapolis St Paul Intl	3870	9E	Endeavor Air Inc.
ATW	Appleton	ORD	Chicago Ohare Intl	5650	OO	SkyWest Airlines Inc.
ATW	Appleton	DTW	Detroit Metro Wayne Co	3694	9E	Endeavor Air Inc.
ATW	Appleton	ORD	Chicago Ohare Intl	5852	EV	ExpressJet Airlines Inc.
ATW	Appleton	MSP	Minneapolis St Paul Intl	4593	OO	SkyWest Airlines Inc.
ATW	Appleton	ORD	Chicago Ohare Intl	6131	EV	ExpressJet Airlines Inc.
ATW	Appleton	MSP	Minneapolis St Paul Intl	3798	9E	Endeavor Air Inc.

3.2.2 LEFT JOIN

While inner joins find intersections, **LEFT JOIN** is exceptionally useful for finding missing or mismatched data.

Let us intentionally look for destination airport codes in the **flights** table that do not have a corresponding entry in the **airports** table. We do this by performing a **LEFT JOIN** and filtering for rows where the airport name resolves to **NULL**.

```

-- using LEFT JOIN

SELECT year, month, day, origin, dest,

```

```

    a.name AS dest_name,
    flight, carrier
FROM flights AS f
LEFT JOIN airports AS a ON f.dest = a.faa
WHERE a.name IS NULL
LIMIT 10;

```

Table 3.9: Displaying records 1 - 10

year	month	day	origin	dest	dest_name	flight	carrier
2013	10	1	JFK	SJU	NA	1403	B6
2013	10	1	JFK	BQN	NA	939	B6
2013	10	1	FLL	SJU	NA	953	B6
2013	10	1	MCO	SJU	NA	933	B6
2013	10	1	MIA	SJU	NA	1341	AA
2013	10	1	SJU	STT	NA	1236	B6
2013	10	1	JFK	SJU	NA	315	DL
2013	10	1	JFK	SJU	NA	3	B6
2013	10	1	EWR	SJU	NA	1289	B6
2013	10	1	EWR	SJU	NA	1216	UA

In applied contexts, queries like this are fundamental to data auditing. In this specific dataset, this query reveals flights traveling to US territories like Puerto Rico (BQN, SJU) or the Virgin Islands (STT). Because they operate slightly outside the standard domestic FAA categorization used to build our `airports` table, they show up as missing values in our join.

i Handling Missing Values in SQL

Users of `R` are used to `NA` and functions like `is.na()`. SQL has a related but stricter concept: `NULL`, representing an **unknown or missing value**. The critical difference is that `NULL` cannot be compared using the standard `=` operator. A test like `a.name = NULL` does not evaluate to `TRUE` or `FALSE` — it evaluates to `UNKNOWN`, and rows where this happens are silently excluded from the result, with no error raised. To correctly test for missing values, `SQL` provides the dedicated keywords `IS NULL` and `IS NOT NULL`. Notice the difference in the `WHERE` clause between the queries above and below..

```
SELECT year, month, day, origin, dest,
       a.name AS dest_name,
       flight, carrier
FROM flights AS f
LEFT JOIN airports AS a ON f.dest = a.faa
WHERE a.name = NULL
LIMIT 10;
```

Table 3.10: 0 records

year	month	day	origin	dest	dest_name	flight	carrier
------	-------	-----	--------	------	-----------	--------	---------

3.3 Unions (Vertical Integration)

While joins add columns to our result set, set operations add rows. `UNION` allows us to append the results of one query directly beneath the results of another, much like `dplyr::bind_rows()`.

Let us extract a dataset combining all flights that flew from ATW to MSP on April 6, 2013, with all flights that flew from JFK to ORD on the same day.

```
-- extract all flights that flew from ATW to MSP and from JFK to ORD on April 6, 2013

(SELECT year, month, day, origin, dest, flight, carrier
FROM flights
WHERE year = 2013 AND month = 4 AND day = 6 AND origin = 'ATW' AND dest = 'MSP')

UNION

(SELECT year, month, day, origin, dest, flight, carrier
FROM flights
WHERE year = 2013 AND month = 4 AND day = 6 AND origin = 'JFK' AND dest = 'ORD')

LIMIT 10;
```

Table 3.11: 8 records

year	month	day	origin	dest	flight	carrier
2013	4	6	ATW	MSP	4559	OO

year	month	day	origin	dest	flight	carrier
2013	4	6	ATW	MSP	3674	9E
2013	4	6	JFK	ORD	905	B6
2013	4	6	JFK	ORD	3521	9E
2013	4	6	JFK	ORD	917	B6
2013	4	6	JFK	ORD	3523	9E
2013	4	6	JFK	ORD	1351	AA
2013	4	6	JFK	ORD	3525	9E

⚠ Union Edge Cases

When using `UNION`, we must be aware of two strict database rules:

- **Column Matching:** Unlike R's `bind_rows()`, which politely aligns columns by their names and fills in NAs where they don't match, `SQL UNION` does not check column names. It blindly stacks the tables based on the number of columns and their data types. If query A returns a string in column 1, and query B returns an integer in column 1, the query will fail.
- **Deduplication:** By default, `UNION` removes exact duplicate rows from the final result set. Searching for and removing duplicates is computationally expensive. If we know our datasets are mutually exclusive (as they are above) or if we want to retain duplicates, we should use `UNION ALL`, which simply appends the rows and executes much faster.

3.4 Subqueries

As our analytical questions grow more sophisticated, we often encounter situations where we cannot get our answer in a single pass. We need to evaluate a query, hold the result in memory, and then query that result.

We do this using subqueries. A **subquery** is simply a complete `SELECT` statement wrapped in parentheses and placed inside the `FROM` or `WHERE` clause of a parent query.

Let us return to the `UNION` query we just wrote above. What if we don't want to see the rows, but rather we just want a count of the total number of flights across both of those specific routes? We can wrap the entire `UNION` in parentheses and run a `COUNT(*)` over it.

```
-- count the number of flights that satisfy the criteria above

SELECT COUNT(*) AS num_flights
```

```

FROM
((SELECT year, month, day, origin, dest, flight, carrier
FROM flights
WHERE year = 2013 AND month = 4 AND day = 6 AND origin = 'ATW' AND dest = 'MSP')
UNION
(SELECT year, month, day, origin, dest, flight, carrier
FROM flights
WHERE year = 2013 AND month = 4 AND day = 6 AND origin = 'JFK' AND dest = 'ORD')
) AS temp_subquery;

```

Table 3.12: 1 records

num_flights
8

The Mandatory Alias

When we place a subquery inside a `FROM` clause, we are essentially creating a temporary table on the fly (often called a **derived table**). In MySQL and MariaDB, every derived table must have its own alias, even if we never explicitly refer to that alias in the `SELECT` statement. Notice the `AS temp_subquery` at the very end of our code block; without it, the SQL engine would throw a syntax error.

3.5 terminate connection

As always, we conclude our workflow by responsibly closing our database connections to free up memory on both our local machine and the remote server.

```

# terminate the SQL connection

DBI::dbDisconnect(con_air, shutdown = TRUE)

```

4 Creating Databases

In our previous chapters, we operated strictly as consumers of data. We connected to a remote, read-only database and executed queries to extract and aggregate information. However, complete data science workflows frequently require us to be data architects. We must ingest raw, unstructured data from flat files, define a rigid structure for it, and build our own relational databases from scratch.

In this chapter, we transition from querying existing systems to architecting our own. To do this, we will use **duckdb**, an in-process SQL database management system. Unlike traditional client-server databases (like MySQL or PostgreSQL) that require external servers and complex authentication, DuckDB runs entirely within our local R environment. It is highly optimized for analytical queries on our local machines, bridging the gap between massive server-side databases and local data frames.

4.1 load packages and connect to database server

Let us load our required packages. Notice the addition of the **duckdb** package alongside our standard database interface tools.

```
library(tidyverse)
library(mdsr)
library(DBI)
library(RMariaDB)
library(dbplyr)
library(duckdb)
```

To contextualize our database creation process, we will consider the Saturday Night Live (SNL) datasets available on the **snldb** GitHub [repo](#). Data is scraped from [SNL archives](#) and [IMDb](#) by Hendrik Hillekes and Colin Morris.

Our first step is to initialize a blank database locally. Rather than pointing to a remote server, we point R to a local directory on our hard drive.

```
# initialize an empty DuckDB database and establish a connection

con_snl <- DBI::dbConnect(duckdb::duckdb(), dbdir = "snl_db")
```

When we run this code, `duckdb` creates a physical file named `snl_db` in our current working directory. Our object, `con_snl`, acts as the connection to this newly minted, empty database.

We can verify that the database is currently empty by running a standard SQL command to show all tables:

```
SHOW TABLES;
```

Table 4.1: 0 records

name

4.2 accessing data (csv files) from GitHub (or, elsewhere)

Before we can build a table in SQL, we must thoroughly understand the shape and nature of the raw data we intend to import. The SNL repository contains several `.csv` files in the **output** directory. Let us examine the `casts` dataset.

We will temporarily read it into R's working memory (see the **Environment** pane) using `readr::read_csv()` strictly for exploratory purposes.

```
# read in the data

casts <- readr::read_csv("https://raw.githubusercontent.com/hhl1cks/snl_db/refs/heads/master/
```

We must inspect the data to identify variable types and the presence of missing values (NAs), as these details dictate how we will construct our SQL schema.

```
glimpse(casts)
```

```
Rows: 614
Columns: 8
$ aid      <chr> "A. Whitney Brown", "A. Whitney Brown", "A. Whitney Br~
$ sid      <dbl> 11, 12, 13, 14, 15, 16, 5, 39, 40, 41, 42, 45, 46, 21, ~
$ featured <lgl> TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, TRUE, ~
```

```

$ first_epid      <dbl> 19860222, NA, NA, NA, NA, NA, 19800409, 20140118, NA, ~
$ last_epid       <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA~
$ update_anchor   <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE~
$ n_episodes      <dbl> 8, 20, 13, 20, 20, 20, 5, 11, 21, 21, 21, 18, 17, 20, ~
$ season_fraction <dbl> 0.4444444, 1.0000000, 1.0000000, 1.0000000, 1.0000000,~

```

The `glimpse()` output reveals how R interprets the data: strings (`chr`), logicals (`lgl`), and numeric values (`dbl`). We also clearly see NA values, which we will need to handle explicitly in SQL.

```
summary(casts)
```

	aid	sid	featured	first_epid
Length	:614	Min. : 1.00	Mode :logical	Min. :19770115
N.unique	:156	1st Qu.:15.00	FALSE:451	1st Qu.:19801215
N.blank	: 0	Median :26.00	TRUE :163	Median :19901110
Min.nchar	: 8	Mean :25.47		Mean :19909634
Max.nchar	: 20	3rd Qu.:37.00		3rd Qu.:19957839
		Max. :46.00		Max. :20141025
				NAs :564
	last_epid	update_anchor	n_episodes	season_fraction
Min.	:19751011	Mode :logical	Min. : 1.00	Min. :0.04167
1st Qu.	:19850112	FALSE:541	1st Qu.:19.00	1st Qu.:1.00000
Median	:19950225	TRUE :73	Median :20.00	Median :1.00000
Mean	:19944038		Mean :18.73	Mean :0.94827
3rd Qu.	:20040117		3rd Qu.:21.00	3rd Qu.:1.00000
Max.	:20140201		Max. :24.00	Max. :1.00000
NAs	:597			

Our summary statistics confirm the presence of missing values. Furthermore, we can deduce that `first_epid` and `last_epid` represent chronological identifiers (episode IDs) rather than standard numeric measurements.

Because SQL databases generally ingest data from local flat files rather than live URLs, our final preparatory step is to save this dataset locally as a `.csv`.

```

# save the dataset as a .csv file on your (local) computer

write_csv(casts, "casts.csv")

```

4.3 data types in SQL

When R imports a `.csv` file, it uses dynamic typing—it scans the first few thousand rows and guesses whether a column is numeric, character, or logical.

SQL, however, requires strict, explicit typing. When we create a database table, we must declare exactly what kind of data each column will hold. This rigorous requirement prevents data corruption (e.g., accidentally inserting the string “five” into a numeric column) and optimizes storage.

Common SQL data types include `VARCHAR` (variable-length string), `INTEGER` (whole numbers), `SMALLINT` (smaller whole numbers to save space), `DOUBLE` (decimals), and `BOOLEAN` (True/False).

Before proceeding, we highly recommend consulting the official [MySQL Data Types Documentation](#) for a comprehensive overview.

4.4 importing data

With our database set up, `.csv` file saved locally, and our data types mapped out, we are ready to build our database table. Importing flat files into a relational database requires a sequential, four-step pipeline:

- a **USE** statement that ensures we are in the right database.
- a series of **DROP TABLE** statements that drop any old tables with the same names as the ones we are going to create.
- a series of **CREATE TABLE** statements that specify the table structures and field types.
- a series of **COPY** (or **LOAD DATA** for MySQL) statements that read the data from the `.csv` files into the appropriate tables.

```
-- step 1: ensure we are in the correct database
```

```
USE snl_db;
```

```
-- step 2: Clear any existing table to ensure reproducibility
```

```
DROP TABLE IF EXISTS casts;
```

Now we define our schema. For every column, we specify the name, the exact data type, whether it is allowed to be empty (NOT NULL), and what the default value should be if a user forgets to provide data. Finally, we define a **Primary Key** — a unique identifier for each row.

```
-- step 3: define the strict schema

CREATE TABLE casts (
  aid VARCHAR(255) NOT NULL DEFAULT ' ',
  sid SMALLINT NOT NULL DEFAULT 0,
  featured BOOLEAN NOT NULL DEFAULT FALSE, /* use TINYINT(1) in MySQL instead of BOOLEAN) */,
  first_epid INTEGER DEFAULT NULL,
  last_epid INTEGER DEFAULT NULL,
  update_anchor BOOLEAN NOT NULL DEFAULT FALSE,
  n_episodes SMALLINT NOT NULL DEFAULT 0,
  season_fraction DOUBLE NOT NULL DEFAULT 0,
  PRIMARY KEY (sid, aid)
);
```

💡 Composite Primary Keys

Notice that our PRIMARY KEY is composed of two columns: (sid, aid). This is known as a **composite key**. It means that while a single actor (aid) might appear multiple times in the table, and a single season (sid) contains multiple actors, the combination of a specific actor in a specific season will only ever appear exactly once.

With our empty structural shell created, we instruct the database engine to ingest the data from our local .csv file.

```
-- step 4: load the data from the flat csv file

COPY casts FROM 'casts.csv' (FORMAT CSV, HEADER, NULL 'NA');
```

4.5 checks

It is vital to audit our work to ensure the data was ingested correctly and the schema holds up.

```
-- verify the database exists

SHOW DATABASES;
```

Table 4.2: 1 records

database_name
snl_db

```
-- verify our table exists

SHOW TABLES;
```

Table 4.3: 1 records

name
casts

```
-- review the strict schema we defined

DESCRIBE casts;
```

Table 4.4: 8 records

column_name	column_type	null	key	default	extra
aid	VARCHAR	NO	PRI	' '	NA
sid	SMALLINT	NO	PRI	0	NA
featured	BOOLEAN	NO	NA	CAST('f' AS BOOLEAN)	NA
first_epid	INTEGER	YES	NA	NULL	NA
last_epid	INTEGER	YES	NA	NULL	NA
update_anchor	BOOLEAN	NO	NA	CAST('f' AS BOOLEAN)	NA
n_episodes	SMALLINT	NO	NA	0	NA
season_fraction	DOUBLE	NO	NA	0	NA

```
-- peek at the ingested data

SELECT * FROM casts LIMIT 10;
```


Table 4.5: Displaying records 1 - 10

aid	sid	fea- tured	first_epid	last_epid	update_an- chor	n_episodes	season_frac- tion
A. Whitney Brown	11	TRUE	19860222	NA	FALSE	8	0.4444444
A. Whitney Brown	12	TRUE	NA	NA	FALSE	20	1.0000000
A. Whitney Brown	13	TRUE	NA	NA	FALSE	13	1.0000000
A. Whitney Brown	14	TRUE	NA	NA	FALSE	20	1.0000000
A. Whitney Brown	15	TRUE	NA	NA	FALSE	20	1.0000000
A. Whitney Brown	16	TRUE	NA	NA	FALSE	20	1.0000000
Alan Zweibel	5	TRUE	19800409	NA	FALSE	5	0.2500000
Sasheer Zamata	39	TRUE	20140118	NA	FALSE	11	0.5238095
Sasheer Zamata	40	TRUE	NA	NA	FALSE	21	1.0000000
Sasheer Zamata	41	FALSE	NA	NA	FALSE	21	1.0000000

4.6 shortcut - a fast (but fragile) route

Many R packages, including `duckdb`, provide convenient wrapper functions that automate the entire pipeline we just built. These functions read the `.csv` file, guess the schema (just like R does), construct the `CREATE TABLE` statement silently, and ingest the data in a single line of code.

Let us demonstrate this shortcut using the `episodes.csv` file.

```
# read in the data

episodes <- readr::read_csv("https://raw.githubusercontent.com/hhlcks/snlldb/refs/heads/master/episodes.csv")
glimpse(episodes)
```

```
Rows: 906
Columns: 4
```

```
$ sid    <dbl> 46, 46, 46, 46, 46, 46, 46, 46, 46, 46, 46, 46, 46, 46, 46, ~
$ epid   <dbl> 20210410, 20210403, 20210327, 20210227, 20210220, 20210213, 2021~
$ aired  <chr> "April 10, 2021", "April 3, 2021", "March 27, 2021", "February 2~
$ epno   <dbl> 17, 16, 15, 14, 13, 12, 11, 10, 9, 8, 7, 6, 5, 4, 3, 2, 1, 18, 1~
```

```
# save the dataset as a .csv file on your (local) computer

write_csv(episodes, "episodes.csv")
```

Now, we can use the `duckdb_read_csv()` function to bypass raw SQL entirely:

```
duckdb::duckdb_read_csv(con = con_sn1, name = "episodes", files = "episodes.csv")
```

We can verify that it worked using SQL:

```
SHOW TABLES;
```

Table 4.6: 2 records

<u>name</u>
casts
<u>episodes</u>

```
DESCRIBE episodes;
```

Table 4.7: 4 records

column_name	column_type	null	key	default	extra
sid	INTEGER	YES	NA	NA	NA
epid	INTEGER	YES	NA	NA	NA
aired	VARCHAR	YES	NA	NA	NA
epno	INTEGER	YES	NA	NA	NA

```
SELECT * FROM episodes LIMIT 10;
```

Table 4.8: Displaying records 1 - 10

sid	epid	aired	epno
46	20210410	April 10, 2021	17
46	20210403	April 3, 2021	16
46	20210327	March 27, 2021	15
46	20210227	February 27, 2021	14
46	20210220	February 20, 2021	13
46	20210213	February 13, 2021	12
46	20210206	February 6, 2021	11
46	20210130	January 30, 2021	10
46	20201219	December 19, 2020	9
46	20201212	December 12, 2020	8

A Fragile Shortcut

While functions like `duckdb_read_csv()` are incredibly useful for rapid prototyping, they are fragile. By allowing the software to guess the schema, we forfeit control over our data architecture. A single malformed string deep in a numeric column could cause the wrapper function to declare the entire column as text, ruining downstream analysis.

In rigorous data engineering contexts, we must demonstrate mastery over data integrity by explicitly writing `CREATE TABLE` schemas. We must dictate the rules, rather than relying on an algorithm to guess them.

4.7 terminate connection

As always, we conclude by cleanly severing our connection to the database.

```
# terminate the SQL connection

DBI::dbDisconnect(con_snl, shutdown = TRUE)
```

Part II

Text Analysis

5 Case Study of Shakespeare's Macbeth using Regular Expressions

5.1 resources

- [Text Mining with R](#)
- [Modern Data Science with R \(3e\) - Chapter 19](#)

5.2 packages

```
library(tidyverse)
```

5.3 data

Our data comes from [Project Gutenberg](#), a large repository of free eBooks. We will be working with the text of Shakespeare's play Macbeth, which is available at <https://www.gutenberg.org/cache/epub/1129/pg1129.txt>. The text is in the public domain, so we are free to use it for our analysis.

```
library(gutenbergr)

data("gutenberg_metadata")

macbeth_all <- gutenberg_metadata |>
  filter(str_detect(title, "Macbeth"))

book <- gutenberg_download(1129, meta_fields = "title")
```

For our case study, we will be working with the raw text of the play, which is available in the `mdsr` package.

```
library(mdsr)

data(Macbeth_raw)

?Macbeth_raw
```

```
macbeth <- Macbeth_raw |>
  str_split("\r\n") |>
  pluck(1)

length(macbeth)

macbeth[300:310]
```

5.4 strings and regular expressions

```
# how many times does the character Macbeth speak in the play?

macbeth_lines <- macbeth |>
  str_subset(" MACBETH")

length(macbeth_lines)

head(macbeth_lines)
```

```
macbeth |>
  str_which(" MACBETH")
```

5.5 about regular expressions grammar

5.5.1 metacharacter

```
macbeth |>
  str_subset("MAC.") |>
  head()
```

```
macbeth |>
  str_subset("MACBETH\\.") |>
  head()
```

5.5.2 character sets

```
macbeth |>
  str_subset("MAC[B-Z]") |>
  head()
```

5.5.3 alternation

```
macbeth |>
  str_subset("MAC(B|D)") |>
  head()
```

5.5.4 anchors (power-dollar)

```
macbeth |>
  str_subset("^ MAC[B-Z]") |>
  head()
```

```
macbeth |>
  str_subset("\\\\.$") |>
  head()
```

5.5.5 repetitions

```
# ?

macbeth |>
  str_subset("^ ?MAC[B-Z]") |>
  head()
```

```
# *

macbeth |>
  str_subset("^ *MAC[B-Z]") |>
  head()
```

```
# +

macbeth |>
  str_subset("^ +MAC[B-Z]") |>
  head()
```

5.6 a case study - life and death in Macbeth and speaking patterns

```
macbeth_chars <- tribble(
  ~name, ~regexp,
  "Macbeth", "  MACBETH\\.",
  "Lady Macbeth", "  LADY MACBETH\\.",
  "Banquo", "  BANQUO\\.",
  "Duncan", "  DUNCAN\\.",
) |>
  mutate(speaks = purrr::map(regexp, str_detect, string = macbeth))
```

```
speaker_freq <- macbeth_chars |>
  unnest(cols = speaks) |>
  mutate(
    line = rep(1:length(macbeth), 4),
    speaks = as.numeric(speaks)
  ) |>
  filter(line > 218 & line < 3172)

glimpse(speaker_freq)
```

```
acts <- tibble(
  line = str_which(macbeth, "^ACT [I|V]+"),
  line_text = str_subset(macbeth, "^ACT [I|V]+"),
  labels = str_extract(line_text, "^ACT [I|V]+")
)
```



```

ggplot(data = speaker_freq, aes(x = line, y = speaks)) +
  geom_smooth(
    aes(color = name), method = "loess",
    se = FALSE, span = 0.4
  ) +
  geom_vline(
    data = acts,
    aes(xintercept = line),
    color = "darkgray", lty = 3
  ) +
  geom_text(
    data = acts,
    aes(y = 0.085, label = labels),
    hjust = "left", color = "darkgray"
  ) +
  ylim(c(0, NA)) +
  xlab("Line Number") +
  ylab("Proportion of Speeches") +
  scale_color_colorblind()

```

6 Case Study of Donald Trump's tweets during the 2016 US presidential election

6.1 resources

- [Trump Twitter Archive](#)
- [Todd Vaziri's Tweet](#)
- [David Robinson's Blog Post](#)
- [Text Mining With R](#)

6.2 packages

```
library(tidyverse)
library(scales)
library(tidytext)
library(textdata)
library(dslabs)
```

6.3 data

Our data come from the *dslabs* package.

```
data("trump_tweets")

?trump_tweets

glimpse(trump_tweets)
```

```
range(trump_tweets$created_at)
```

```
trump_tweets$text[16413]
```

```
trump_tweets |>  
  count(source) |>  
  arrange(desc(n))
```

For our case study, we are interested in what happened during the 2016 campaign, so for this analysis we will focus on what was tweeted between the day Trump announced his campaign and election day. Also, we only keep tweets from Android or iPhone, and filter out retweets.

```
campaign_tweets <- trump_tweets |>  
  filter(source %in% paste("Twitter for", c("Android", "iPhone"))) &  
    created_at >= ymd("2015-06-17") &  
    created_at < ymd("2016-11-08")) |>  
  mutate(source = str_remove(source, pattern = "Twitter for ")) |>  
  filter(!is_retweet) |>  
  arrange(created_at) |>  
  as_tibble()
```

```
campaign_tweets |>  
  count(source)
```

6.4 some initial exploration

We explore the possibility that two different groups were tweeting from these devices.

```
campaign_tweets |>  
  mutate(hour = hour(with_tz(created_at, "EST"))) |>  
  count(source, hour)
```

```
android_vs_iphone_hourly_props <- campaign_tweets |>  
  mutate(hour = hour(with_tz(created_at, "EST"))) |>  
  count(source, hour) |>  
  group_by(source) |>  
  mutate(percent = n / sum(n))  
  
android_vs_iphone_hourly_props
```

```
android_vs_iphone_hourly_props |>
  ggplot(aes(x = hour, y = percent, color = source)) +
  geom_line() +
  geom_point() +
  scale_y_continuous(labels = percent_format()) +
  labs(x = "Hour of day (EST)", y = "% of tweets", color = "")
```

6.5 tokens

```
i <- 3008

campaign_tweets$text[i]
```

```
campaign_tweets[i,] |>
  tidytext::unnest_tokens(output = word, input = text, token = "words") |>
  dplyr::pull(word)
```

```
links_to_pics <- "(https://t.co/[A-Za-z\\d]+)|(http://t.co/[A-Za-z\\d]+)"

campaign_tweets[i,] |>
  mutate(text = str_remove_all(text, pattern = links_to_pics)) |>
  tidytext::unnest_tokens(output = word, input = text, token = "words") |>
  dplyr::pull(word)
```

```
tweet_words <- campaign_tweets |>
  mutate(text = str_remove_all(text, pattern = links_to_pics)) |>
  tidytext::unnest_tokens(output = word, input = text, token = "words")
```

```
tweet_words |>
  count(word) |>
  arrange(desc(n))
```

6.6 stop words

```
data(stop_words)
```

```
?stop_words
```

```
tweet_words <- campaign_tweets |>
  mutate(text = str_remove_all(text, pattern = links_to_pics)) |>
  tidytext::unnest_tokens(output = word, input = text, token = "words") |>
  filter(!(word %in% stop_words$word))
```

```
tweet_words |>
  count(word)
```

```
tweet_words <- campaign_tweets |>
  mutate(text = str_remove_all(text, links_to_pics)) |>
  tidytext::unnest_tokens(output = word, input = text, token = "words") |>
  filter(!(word %in% stop_words$word)) |>
  filter(!(str_detect(word, pattern = "^\\d+")))
```

```
tweet_words |>
  count(word)
```

```
# |>
#   arrange(desc(n))
```

6.7 further exploration

```
android_vs_iphone <- tweet_words |>
  count(word, source)
```

```
android_vs_iphone
```

```
android_vs_iphone <- tweet_words |>
  count(word, source) |>
  pivot_wider(names_from = "source", values_from = "n", values_fill = 0)
```

```
android_vs_iphone
```

```

android_vs_iphone <- tweet_words |>
  count(word, source) |>
  pivot_wider(names_from = "source", values_from = "n", values_fill = 0) |>
  mutate(p_a = (Android + 0.5)/(sum(Android) + 0.5),
         p_i = (iPhone + 0.5)/(sum(iPhone) + 0.5),
         ratio = p_a / p_i)

android_vs_iphone

```

for words appearing at least 100 times in total, here are the ten highest percent differences

```

top_10_android <- android_vs_iphone |>
  filter(Android + iPhone >= 100) |>
  arrange(desc(ratio)) |>
  slice_head(n = 10)

top_10_android

```

and the top ten for iPhone

```

top_10_iphone <- android_vs_iphone |>
  filter(Android + iPhone >= 100) |>
  arrange(ratio) |>
  slice_head(n = 10)

top_10_iphone

```

```

top_10_android |>
  rbind(top_10_iphone) |>
  dplyr::select(word, iPhone, Android, ratio) |>
  mutate(log_ratio = log(ratio)) |>
  mutate(device = rep(c("Android", "iPhone"), each = 10)) |>
  ggplot() +
  geom_col(aes(x = log_ratio, y = fct_reorder(word, log_ratio), fill = device))

```

6.8 sentiment analysis

```

tidytext::get_sentiments("bing")

```

```
tidytext::get_sentiments("afinn")
```

```
tidytext::get_sentiments("nrc")
```

```
nrc <- tidytext::get_sentiments("nrc")
```

```
nrc |> count(sentiment)
```

```
tweet_words |>  
  left_join(nrc, by = "word") |>  
  dplyr::select(source, word, sentiment) |>  
  sample_n(5)
```

```
sentiment_counts <- tweet_words |>  
  left_join(nrc, by = "word", relationship = "many-to-many") |>  
  count(source, sentiment)
```

```
sentiment_counts
```

```
sentiment_counts <- tweet_words |>  
  left_join(nrc, by = "word", relationship = "many-to-many") |>  
  count(source, sentiment) |>  
  pivot_wider(names_from = "source", values_from = "n") |>  
  mutate(sentiment = replace_na(sentiment, replace = "none"))
```

```
sentiment_counts
```

```
sentiment_counts <- sentiment_counts |>  
  mutate(p_a = Android/sum(Android),  
         p_i = iPhone/sum(iPhone),  
         ratio = p_a/p_i) |>  
  arrange(desc(ratio))
```

```
sentiment_counts
```

```
ordered_levels <- sentiment_counts$sentiment
```

```
android_vs_iphone |>  
  filter(Android + iPhone >= 10) |>  
  left_join(nrc, by = "word") |>
```

```
mutate(sentiment = factor(sentiment, levels = sentiment_counts$sentiment)) |>
group_by(sentiment) |>
slice_max(abs(log(ratio)), n = 10) |>
ungroup() |>
mutate(word = reorder(word, -ratio)) |>
ggplot(aes(word, ratio, fill = ratio < 1)) +
geom_col(show.legend = FALSE) +
scale_y_log10() +
facet_wrap(~sentiment, scales = "free_x", nrow = 2) +
theme(axis.text.x = element_text(angle = 90, hjust = 1))
```


7 Case Study of Data Science Papers from arXiv.org

7.1 resources

- [arXiv.org](https://arxiv.org)

7.2 packages

```
library(tidyverse)
library(scales)
library(tidytext)
library(textdata)
library(mdsr)
```

7.3 data

We can import data using the **aRxiv** package.

```
library(aRxiv)

df <- arxiv_search(
  query = 'ti:"Data Science"',
  limit = 20000,
  batchsize = 100
)

glimpse(df)
```

For this case study we will use some already-collected data from the **mdsr** package.

```
library(mdsr)

data(DataSciencePapers)

?DataSciencePapers

glimpse(DataSciencePapers)
```

7.4 initial data exploration and cleaning

```
DataSciencePapers <- DataSciencePapers |>
  mutate(
    submitted = lubridate::ymd_hms(submitted),
    updated = lubridate::ymd_hms(updated)
  )

glimpse(DataSciencePapers)
```

```
mosaic::tally(~ year(submitted), data = DataSciencePapers)
```

```
DataSciencePapers |>
  filter(id == "1809.02408v2") |>
  glimpse()
```

```
DataSciencePapers |>
  group_by(primary_category) |>
  count()
```

```
DataSciencePapers <- DataSciencePapers |>
  mutate(field = str_extract(primary_category, pattern = "[a-z,-]+"))

mosaic::tally(x = ~field, margins = TRUE, data = DataSciencePapers) |>
  sort()
```

7.5 corpora and tokenization

```
DataSciencePapers |>
  unnest_tokens(output = word, input = abstract, token = "words") |>
  count(id, word, sort = TRUE)
```

```
arxiv_words <- DataSciencePapers |>
  unnest_tokens(output = word, input = abstract, token = "words") |>
  anti_join(tidytext::get_stopwords(), by = "word")

arxiv_words |>
  count(id, word, sort = TRUE)
```

```
arxiv_abstracts <- arxiv_words |>
  group_by(id) |>
  summarize(abstract_clean = paste(word, collapse = " "))

arxiv_papers <- DataSciencePapers |>
  left_join(arxiv_abstracts, by = "id")
```

7.6 wordcloud

```
library(wordcloud)

set.seed(4302026)
arxiv_papers |>
  pull(abstract_clean) |>
  wordcloud(
    max.words = 40,
    scale = c(8, 1),
    colors = topo.colors(n = 30),
    random.color = TRUE
  )
```

7.7 bigrams

```
arxiv_bigrams <- arxiv_papers |>
  unnest_tokens(
```

```

    output = arxiv_bigram,
    input = abstract_clean,
    token = "ngrams",
    n = 2
  ) |>
  dplyr::select(arxiv_bigram, id)

arxiv_bigrams

arxiv_bigrams |>
  count(arxiv_bigram, sort = TRUE)

```

7.8 tf-idf and document term matrices

```

arxiv_words |>
  count(word) |>
  arrange(desc(n)) |>
  head()

```

```

tidy_DTM <- arxiv_words |>
  count(id, word) |>
  tidytext::bind_tf_idf(word, id, n)

tidy_DTM

```

```

tidy_DTM |>
  arrange(desc(tf)) |>
  head()

```

```

tidy_DTM |>
  arrange(desc(idf), desc(n)) |>
  head()

```

```

tidy_DTM |>
  filter(word == "wildfire")

```

```

arxiv_papers |>
  pull(abstract) |>

```

```
str_subset("wildfire") |>  
strwrap() |>  
head()
```

```
tidy_DTM |>  
  filter(word == "implications")
```

```
tidy_DTM |>  
  filter(id == "1809.02408v2") |>  
  arrange(desc(tf_idf)) |>  
  head()
```

```
tm_DTM <- arxiv_words |>  
  count(id, word) |>  
  tidytext::cast_dtm(document = id, term = word, value = n)  
  
tm_DTM
```

Part III

Geospatial Data

8 Recreating John Snow's 1854 map of the Broad Street cholera outbreak.

8.1 resources

- [Modern Data Science with R \(3e\) - Chapter 17](#)
- [1854 Broad Street cholera outbreak](#)
- [Shapefile](#)

8.2 packages

```
library(tidyverse)
library(sf)
library(ggspatial)
library(maps)
library(mapproj)
```

8.3 data

```
root <- here::here("/cloud/project") # where we stored these files
dsn <- fs::path(root, "SnowGIS_SHP")
```

```
list.files(dsn)
```

```
st_layers(dsn)
```

```
CholeraDeaths <- st_read(dsn, layer = "Cholera_Deaths")
```

```
class(CholeraDeaths)
```

```
plot(CholeraDeaths["Count"])
```

```
ggplot(CholeraDeaths) +  
  geom_sf()
```

```
ggplot(CholeraDeaths) +  
  ggspatial::annotation_map_tile(type = "osm", zoomin = 0, progress = "none") +  
  geom_sf(aes(size = Count), alpha = 0.7)
```

```
# bounding boxes
```

```
st_bbox(CholeraDeaths)
```

8.4 projections

[Top 10 World Map Projections](#)

```
maps::map(database = "world", projection = "mercator", wrap = TRUE)
```

```
maps::map(database = "world", projection = "cylequalarea", param = 45, wrap = TRUE)
```

8.5 CRS - coordinate reference system

[Coordinate Reference System](#) - three formats that are common for storing information about the projection of a geospatial object are EPSG, PROJ.4, and WKT

Common CRSs

- [EPSG:4326](#) - also known as WGS84, standard for GPS systems and Google Earth
- [EPSG:3857](#) - a Mercator projection used in map tiles by Google Maps, Open Street Maps, etc.
- [EPSG:27700](#) - also known as OSGB 1936, or the British National Grid, commonly used in Britain


```
st_crs(4326)$epsg  
  
st_crs(3857)$epsg  
  
st_crs(27700)$epsg
```

```
st_crs(4326)$Wkt  
  
st_crs(3857)$Wkt  
  
st_crs(27700)$Wkt
```

```
st_crs(4326)$proj4string  
  
st_crs(3857)$proj4string  
  
st_crs(27700)$proj4string
```

8.6 what about our dataset?

```
st_crs(CholeraDeaths)
```

```
cholera_4326 <- CholeraDeaths |>  
  st_transform(4326)  
  
st_bbox(cholera_4326)
```

```
ggplot(cholera_4326) +  
  annotation_map_tile(type = "osm", zoomin = 0) +  
  geom_sf(aes(size = Count), alpha = 0.7)
```

```
# add pump locations  
  
pumps <- st_read(dsn, layer = "Pumps")  
  
pumps_4326 <- pumps |>  
  st_transform(4326)
```

```
ggplot(cholera_4326) +  
  annotation_map_tile(type = "osm", zoomin = 0) +  
  geom_sf(aes(size = Count), alpha = 0.7) +  
  geom_sf(data = pumps_4326, color = "red", size = 3)
```

Part IV

Code Performance

9 Code Performance

9.1 resources

- [Advanced R](#) - Chapters 22-25

9.2 code profiling

```
# load the package
library(profvis) # for profiling R code and visualizing the results
```

```
# example 1

f <- function() {
  pause(0.1)
  g()
  h()
}

g <- function() {
  pause(0.1)
  h()
}

h <- function() {
  pause(0.1)
}

####

profvis(f())
```

```
# example 2

# a slow, iterative random walk in R
random_walk_r <- function(n) {
  x <- numeric(n)
  x[1] <- 0
  for(i in 2:n) {
    x[i] <- x[i-1] + rnorm(1, mean = 0, sd = 1)
  }
  return(x)
}

####

profvis::profvis({
  # running a long simulation to give the profiler enough data
  sim_data <- random_walk_r(1e5)
})
```

9.3 improving performance with Rcpp

```
# load the package
library(Rcpp) # for integrating C++ code into R

# define the C++ function as a string (or in a separate .cpp file)

Rcpp::cppFunction('
  NumericVector random_walk_cpp(int n) {
    NumericVector x(n);
    x[0] = 0;
    for(int i = 1; i < n; ++i) {
      x[i] = x[i-1] + R::rnorm(0, 1);
    }
    return x;
  }
')
```

OR see below

```
#include <Rcpp.h>
using namespace Rcpp;

// [[Rcpp::export]]

NumericVector random_walk_cpp_2(int n) {
  NumericVector x(n);
  x[0] = 0;
  for(int i = 1; i < n; ++i) {
    x[i] = x[i-1] + R::rnorm(0, 1);
  }
  return x;
}
```

```
profvis::profvis({
  # running a long simulation to give the profiler enough data
  sim_data_cpp <- random_walk_cpp(1e5)
})
```

```
# install.packages("bench")
library(bench)

# compare the R vs C++ implementations
results <- bench::mark(
  r_version = random_walk_r(10000),
  cpp_version = random_walk_cpp(10000),
  iterations = 100,
  check = FALSE # don't check for exact identical output due to random seeds
)
```

```
results |>
  dplyr::select(expression, min, median, `itr/sec`, mem_alloc)
```

```
plot(results)
```

9.4 parallelization with future and furrr

```
# load the packages
library(tictoc) # for timing our code execution
```

```
library(future)    # the core parallel backend architecture
library(furrr)     # a parallel implementation of the purrr family of functions
```

```
future::availableCores()
```

```
future::plan(multisession, workers = 2)
```

```
simulate_median <- function(iteration, n = 1000) {
  # simulate heavily skewed data (e.g., Log-Normal)
  sample_data <- rlnorm(n, meanlog = 10, sdlog = 1)

  # return a tibble with the result
  tibble(
    iteration = iteration,
    median_val = median(sample_data)
  )
}
```

```
# always test the function at least once before parallelizing
```

```
simulate_median(iteration = 1)
```

```
tic("Sequential Purrr")
set.seed(405)
results_seq <- 1:1e4 |>
  purrr::map_dfr(.f = simulate_median)
toc()
```

```
tic("Parallel Furrr")
set.seed(405)
results_par <- 1:1e4 |>
  furrr::future_map_dfr(.f = simulate_median, .options = furrr_options(seed = TRUE))
toc()
```

9.5 efficient file handling

```
# load the packages
library(tidyverse)    # for data manipulation and writing CSV files
library(nanoparquet)  # for writing Parquet files in R
```

```
write_csv(results_seq, "results_seq.csv")

write_parquet(results_seq, "results_seq.parquet")

read_bench <- bench::mark(
  csv = read_csv("results_seq.csv", show_col_types = FALSE),
  parquet = read_parquet("results_seq.parquet"),
  iterations = 15,
  check = FALSE
)

read_bench |> select(expression, median)

plot(read_bench)
```


Part V

R Package

10 Building an R Package

10.1 check for available package names

```
# check available package names (chapter 4.1.2 of R Packages (2e))
install.packages(c("available", "pak"))
available("package_name")
pak::pkg_name_check("package_name")
```

10.2 create package

```
# start by creating a GitHub repo with the package name
# clone that repo to your Posit Cloud
```

```
# install some packages which are required for package development
install.packages(c("devtools", "checkmate", "tidyverse", "remotes"))
# devtools - meta package for package development in R
# checkmate - to build assertions (make sure that our functions are working fine)
# tidyverse - to build one of our functions
# remotes is a package that is required to install a package
```

```
library(devtools)
```

```
# create the package files (basic)
usethis::create_package("/cloud/project")
# make sure to hit the 'Yes' option and check the files/folders that have been created
# add your package name to the Package field in the DESCRIPTION file
# commit and push to GitHub
```

```
library(devtools)
```

10.3 create functions within the package

```
# existing function called strsplit that comes with base R
strsplit("alpha, bravo", split = ",")
# this returns by default a list which is annoying
strsplit("alpha, bravo", split = ",")[[1]]
# we will create a new function called strsplit1 that returns a vector instead of a list
```

```
# make sure the directory is okay
getwd() # should see /cloud/project
```

```
# every function in an R package needs to be written in an .R script with the same name
# which will be stored in the R directory
usethis::use_r("strsplit1")
# type in function R script
strsplit1 <- function(x, split)
{
  strsplit(x, split = split)[[1]]
}
```

```
devtools::load_all()
# loads the package from source directly into the current environment
# temporarily builds and installs the package, allows us to verify our functions
```

```
# check if your functions works
strsplit1("alpha, bravo", split = ",")
# commit and push to GitHub
```

```
# create the min_max_scale function
usethis::use_r("min_max_scale")
# type in function R script
min_max_scale <- function(x)
{
  (x - min(x, na.rm = TRUE)) / (max(x, na.rm = TRUE) - min(x, na.rm = TRUE))
}

devtools::load_all()

# check
min_max_scale(1:5)
# commit and push to GitHub
```

```

# create a personal theme (theme_academic)
usethis::use_r("theme_academic")
# type in function R script
theme_academic <- function(base_size = 12) {
  ggplot2::theme_classic(base_size = base_size) +
    ggplot2::theme(
      plot.title = ggplot2::element_text(face = "bold", hjust = 0.5),
      panel.grid.major = ggplot2::element_blank(),
      panel.grid.minor = ggplot2::element_blank()
    )
}
# specifically mention what external packages are used to build functions in your package
usethis::use_package("ggplot2")
# check Imports field in DESCRIPTION file

devtools::load_all()

# check
library(ggplot2)
p <- ggplot(mtcars) + geom_point(aes(x = mpg, y = hp)) + labs(title = "plot title")
p
p + theme_academic()
# commit and push to GitHub

```

```

# so we have created the barebones version of all our three functions

```

```

# we updated the DESCRIPTION file (changes to title, authors, description)
# most of the other things will get updated on their own
devtools::check()

# run the above check almost after every step
# should get 0 errors, 0 warnings, 0 notes
# right now we get a warning/note about license

```

10.4 add license

```

# let's add a license
usethis::use_mit_license()
# we won't change the LICENSE and LICENSE.md file by hand

```

```
devtools::check() # check again
# commit and push to GitHub
```

10.5 add documentations

```
# we are now going add documentations to our functions
# go to the function .R script > Code> Insert roxygen skeleton
# make sure to have the @export tag
# think of a user when you are writing documentations
devtools::document() # this needs to be run everytime we update the documentation
devtools::load_all()
devtools::check()
# commit and push to GitHub
```

10.6 add assertions

```
# we will work with assertions from the checkmate package
# first check of the function arguments so that it runs smoothly
# checkmate provides somewhat better messages
usethis::use_package("checkmate")
# we will write assertions for each input of each function
# see help pages of ?assert_
devtools::load_all()
# add assertions to each function and run devtools::load_all() after each change
# tryto break your function and see if the assertion works
devtools::document()
devtools::check()
# commit and push to GitHub
```

10.7 unit testing

```
# unit testing of functions
# think about every possible sceanrio where your function can break
usethis::use_testthat()
```

```

usethis::use_test("strsplit1")
# generate a new file inside the testthat directory inside tests directory
# we will delete the prepopulated message and write our own for our functions
# see help pages of ?expect_
# write your unit tests
devtools::load_all()
devtools::test()

```

```

usethis::use_test("min_max_scale")
devtools::load_all()
devtools::test()

```

```

usethis::use_test("theme_academic")
devtools::load_all()
devtools::test()

```

```

devtools::check()
# commit and push to GitHub

```

10.8 add dataset

```

# add dataset to our package that can be exported and used by our users
# first make sure that the dataset is created in your local Environment
toy_data <- data.frame(
  subject_id = 1:5,
  reaction_time = c(1.2, 0.9, 1.5, 0.8, 1.1),
  test_score = c(45, 82, 55, 91, 67),
  recalled_words = c(
    "cat, dog, mouse",
    "apple, banana",
    "car, truck, bike, train",
    "red, blue, green",
    "sun, moon"
  )
)

```

```

usethis::use_data(toy_data)
# if exporting big datasets, look into the compress argument of use_this

```

```
# we will save the data generating script into a file (for reproducibility)
# very important if you generate datasets that need random generation (set a seed)
usethis::use_data_raw("toy_data_code")
# anytime you update the toy_data with new code, make sure to update the toy_data_code file
```

```
# documentation for toy_data
usethis::use_r("toy_data")
# seems to be done manually
# this will build the documentation
# make sure that the last line has the name of the dataset inside ""
# mention the units of your variables
```

```
# if we export multiple datasets, R puts them all in sysdata.rda file
# in the NAMESPACE, it doesn't specifically mention the dataset to be exported
```

```
devtools::document()
devtools::load_all() # display the dataset and check help page
devtools::check()
# commit and push to GitHub
```

```
# for any file that is typically not an rda file, like csv files
# we will need to create a directory called inst/extdata
```

10.9 add README

```
# we will build the README for the GitHub site and also for our website
usethis::use_readme_rmd()
# try not to install any other packages (since those packages have not been mentioned in Imports)
devtools::build_readme()
# use this build function repeatedly whenever you update the README.Rmd
devtools::check()
# commit and push to GitHub
# check GitHub
```

10.10 add/build vignettes

```
# build a vignette
# these are specific instructions/notes about very specific part of your package
# for example check out the ggplot2 vignettes from Articles dropdown of the package website
usethis::use_vignette(name = "intro-to-toypackage", title = "Introduction to toypackage")
devtools::build_vignettes() # to convert the Rmd into website readable html
devtools::check()
# commit and push to GitHub
```

10.11 build package website

```
# build the package website
# pkgdown package this is how the website will be built
usethis::use_pkgdown() # run this only once
# the website will also use GitHub pages like it did for our personal website
pkgdown::build_site() # run this everytime there is an update to anything inside your package
usethis::use_pkgdown_github_pages() # notice the changes, run this only once
devtools::check()
# commit and push to GitHub
```